

When Are Agent Execution Edits Safe?

An Exact Theory of Fork, Restore, and Merge

Abstract—Agent runtimes can Fork an execution, Restore a checkpoint, or Merge branches without restarting a task. We call these operations *execution edits* because they change which parts of the workflow will run next. An edit cannot undo a tool call that has already been authorized or released. A careless edit can therefore authorize or release the same protected action twice, drop work that the workflow still requires, or overlap with a call that began before the edit. We ask when an execution edit can be implemented safely. Our answer preserves every outcome that was required before the edit. Each such outcome must remain required, already be satisfied by the recorded execution, or be explicitly removed by the same authority that required it. Existing Agent recovery systems implement particular edits, while existing synthesis methods assume that the possible executions and completion conditions are already known. Neither derives the full safety requirement of a requested edit from the execution record. We present an exact theory and an executable checker. Given a registered workflow, its authenticated record, and a requested edit, the checker constructs every complete execution allowed by the workflow and policy. It removes an execution whenever one of its partial executions would leave a still-required outcome with no safe way to finish. If nothing remains, the checker returns a finite, independently checkable explanation that no safe implementation exists. Otherwise, the remaining executions define the most permissive runtime monitor that enforces the edit. For all six registered execution edits and registered extensions, our main theorem proves that the checker accepts exactly when such a safe implementation exists. The same theorem proves atomic installation despite racing calls, identifies the workflow, call identity, prior authorization, and monitor-version information that no exact checker can discard, and relates the installed runtime to an ideal atomic machine. Every supported event preserves the security invariant, so the theorem applies again after every finite reachable execution. Six Lean modules mechanize registration, all six edit rules, atomic monitor installation, call-installation races, and preservation across finite executions; 19 paper-specific tests and measurements over 2–128 outcomes exercise the executable checker.

I. INTRODUCTION

Tool-using Agent runtimes no longer follow only one linear plan. They let an Agent Fork an execution into alternatives, Restore an earlier checkpoint, and Merge branches [1], [2], [3], [4], [5], [6]. These operations support exploration, recovery, and reuse without restarting the task. We call them *execution edits* because they change which parts of the workflow will run next.

An execution edit does not undo earlier tool calls. Suppose an Agent saves checkpoint k before a protected approval. The approval is later authorized once, creating an authorization record that Restore does not delete. Restoring k creates another approval call in the workflow, but it must not authorize the same approval again. A concurrent Merge can cause a similar

error by selecting supplier L after a supplier- R payment has already been authorized. An edit can therefore restore the saved program state yet violate authorization or release order; the remote service reports physical completion separately through an authenticated completion record.

The problem is that a checkpoint contains only part of the information needed to decide whether an edit is safe. The controller must also know four things: which outcomes remain required, which call positions refer to the same protected action, which actions have already been authorized, and which runtime version governs a call that overlaps the edit. Each can change the correct answer. The controller cannot recover all four from equal workspaces, an execution trace alone, an authorization log alone, or a replacement workflow proposed by the Agent.

Existing approaches cover parts of this problem but do not derive its complete safety condition. Agent systems provide branching, staged effects, transactions, or recovery checks [7], [8], [9], [10], [11], [12], [13], [14]. Control and synthesis choose behavior after a model of possible executions and desired final states has been supplied [15], [16], [17], [18]. Neither line of work derives from a recorded execution everything that a particular edit must preserve before it takes effect.

Our key idea is that an execution edit must preserve both past actions and still-required outcomes. Every outcome required before the edit must remain required, already be satisfied by the recorded execution, or be explicitly removed by the same authority that required it. The Agent selects only a registered edit. The platform authenticates the registered workflow, and the trusted controller derives the edited workflow and inherited requirements from that record rather than from Agent text. The Agent therefore cannot obtain a favorable answer by submitting a replacement workflow that hides inconvenient facts.

This idea creates three technical problems. First, calls copied by Fork or Restore may refer to the same protected action, so the controller must distinguish a new action from reuse of an earlier call and return value. Second, after every allowed partial execution, every outcome that is still required and still possible must retain a safe way to finish. The checker must enforce this condition without forbidding additional safe behavior. Third, the answer must take effect atomically with calls that overlap the edit and remain valid if the runtime or machine stops unexpectedly and later recovers.

We give an exact checker for execution edits: it accepts exactly when a safe implementation exists under the model defined in section III. For a requested edit, the controller uses the registered workflow and current authenticated state to

construct every allowed completed execution. It removes an execution whenever one of its partial executions would leave a still-required outcome with no safe completion. If no execution remains, the controller returns an independently checkable rejection explanation. Otherwise, the remaining executions define the most permissive safe monitor. Before the edit takes effect, the runtime checks the current authenticated state again and atomically replaces the old monitor. Every overlapping call is therefore governed entirely by either the old monitor or the new one. Section II derives this process for a protected approval workflow containing Fork, Merge, both installation races, and Restore. After Restore, the repeated approval call reuses the earlier authorization record instead of authorizing the approval again.

a) Contributions: We contribute one end-to-end theorem connecting the original Agent workflow, the requested execution edit, the exact checker, and the installed runtime. For all six registered execution edits and registered extensions, a safe implementation exists exactly when the checker leaves a nonempty set of executions, its monitor can be installed, and an ideal atomic machine permits the same edit. The theorem also proves that an exact checker cannot discard workflow structure, the mapping from call positions to protected actions, prior authorization state, or the current monitor version, and that the concrete runtime has the same visible behavior as the ideal machine (sections III to VI). Every supported event preserves the security invariant, so the theorem applies again after every finite reachable execution. Six Lean modules mechanize the central theorem chain; 19 paper-specific executable tests and measurements over 2–128 outcomes exercise the checker and edit-installation races.

II. EXECUTION-EDIT SAFETY

Consider an Agent operating an approval workflow with one authorized order. The Agent reserves the budget and saves checkpoint k before approval. The approval is then authorized once, creating an authorization record that Restore does not delete. The Agent next requests a Fork between suppliers L and R . The workflow allows two outcomes: order from L or order from R . In either outcome, payment authorization must precede shipment authorization:

$$\begin{aligned} o_L : x_L^p < x_L^s, \\ o_R : x_R^p < x_R^s, \end{aligned} \quad (1)$$

where x_i^p and x_i^s are payment and shipment occurrences. An occurrence identifies one call position in the workflow. An effect identity d identifies the protected action and one-use permission behind that position.

The two shipment calls share one effect identity:

$$q_{\text{eff}}(x_L^s) = q_{\text{eff}}(x_R^s) = d_s.$$

The two payment calls instead have different effect identities d_L^p and d_R^p .

The authorization log Δ records newly authorized effects in order. The notation $\text{lab}(\Delta)$ keeps only their permission labels. Budget reservation contributes r , approval contributes a , the

two payments contribute p_L and p_R , and shipment contributes s . At the Fork request, $\text{lab}(\Delta) = ra$. The policy allows prefixes of rap_Ls and rap_Rs , but never both payments or a payment before approval. Although the supplier outcomes are exclusive, both must still have a safe way to finish when the Fork is checked.

A. Accepted and Rejected Forks

The Agent submits only $u = \text{ForkChoice}(b, \sigma)$. The trusted controller, not the Agent, constructs the edited workflow from the registered edit schema. From the authenticated workflow, execution record, authorization log, and policy, it derives the edited workflow, its required outcomes, and the effect identity behind every call position.

Write M_o for the policy-safe completions of outcome o . Each outcome has one allowed call order:

$$\begin{aligned} z_L &= \text{new}(x_L^p, d_L^p) \text{new}(x_L^s, d_s), \\ z_R &= \text{new}(x_R^p, d_R^p) \text{new}(x_R^s, d_s). \end{aligned}$$

Thus $M_{o_L} = \{z_L\}$ and $M_{o_R} = \{z_R\}$. Each supplier outcome is preserved, and every permitted prefix can still finish safely. The verdict is therefore *Accept*: both supplier outcomes remain possible, every permitted prefix is policy safe, and each such prefix can still reach a complete outcome. If policy excludes rap_Rs , then $M_{o_R} = \emptyset$, and the Fork is rejected because pruning supplier R would change its meaning. Rejection returns a certificate listing the order in which completions were removed and the reason for each removal, rather than letting the Agent weaken the requested edit. The verdict in this variant is *Reject*: checking only the surviving L branch would silently weaken the requested Fork.

B. Why Safe Endpoints Are Not Enough

The previous rejection is visible at the start because one outcome has no safe completion. A harder failure occurs when every outcome is initially safe in isolation, but no single monitor can keep all of them completable. Let X, Y, Z be single-call contracts with permission labels x, y, z , respectively, and consider

$$X \otimes (Y \oplus Z) \quad \text{under} \quad \mathcal{A} = \text{Pref}\{yx, xz\}.$$

The two outcomes have safe complete traces yx and xz . However, executing the shared x call first keeps the $X \otimes Y$ outcome possible while leaving only the forbidden continuation xy . A monitor that exposes x is therefore unsafe, while a monitor that removes x destroys the $X \otimes Z$ outcome. The exact verdict is *Reject*, even though an outcome-by-outcome check would accept. Section IV detects this failure by repeatedly removing executions that create such a bad prefix.

C. Installation Races

Figure 1 separates the active workflow, the authorization log, and the monitor version that must change atomically.




Fig. 1. Fork, Restore, and Merge change the active workflow but do not erase recorded calls. Restore creates a new call position while preserving its authenticated effect identity. The atomic installation point orders both creation and reuse of an authorization record against the monitor-version change.

a) *New authorization during installation:* After the shared approval and before either payment, the Agent requests $\text{MergeSelect}(g, L, \sigma)$. If the payment call x_R^p from the old supplier- R branch is authorized first, it selects R and changes $\text{lab}(\Delta)$ from ra to rap_R . Selecting L must therefore be re-derived and rejected. If the Merge installation occurs first, it closes the old monitor version and gives every current call a new authorization token before any later x_R^p can be authorized. Leaving the old monitor version active would allow both payments.

b) *Authorization reuse during installation:* Now consider the Restore that keeps both branches in figure 1. A later Restore of k creates a new approval occurrence x_a' , but both approval occurrences refer to the same effect identity d_a . The restored occurrence is a new workflow position but retains effect identity d_a . Because d_a already has an authorization record, executing x_a' is an reuse: the restored branch advances past approval without creating another authorization record. A concurrent Merge or Restore result computed before that progress is now out of date even though the authorization log has not changed. The installation check must therefore compare both the authorization log and the ordered record χ of executed call positions. For this RestoreLive request, the checker either installs the derived monitor under a new version or returns a rejection certificate.

These executions determine the three requirements used throughout the paper. First, an edit must preserve every outcome that the checked state still requires. Second, after every permitted prefix, every outcome still compatible with that prefix must retain a safe completion. Third, the derived monitor and its authorization tokens must be installed atomically against concurrent protected calls. Section III now defines the authenticated state and fixes this security objective before section IV constructs the checker.

III. MODEL AND SECURITY OBJECTIVE

The model separates unrestricted internal Agent reasoning from tool calls whose effects pass through a trusted runtime monitor. Such a call must identify its registered position in the workflow, present an authorization token, and submit a concrete tool request. The monitor decides whether the call

may use a new authorization, reuse an authorization already recorded in the log, or make no progress. Each modeled state covers one *policy domain*: the protected actions governed by one policy, one authorization log, and one atomically replaced monitor. Domains with disjoint protected actions, authorization logs, and policies compose by product. Every state snapshot is finite, although an execution may contain an unbounded number of snapshots. The definitions below first describe the required outcomes and then the authenticated record needed to preserve them across edits.

A. Workflow outcomes and call order

a) *Semantic vocabulary:* We capitalize *Agent* for the untrusted principal requesting execution edits. An *occurrence* x identifies one position for a protected tool call in the workflow. An *effect identity* d identifies the one protected action and one-use permission behind that position. The first authorized use of an effect identity creates an authorization record. After Fork or Restore, several occurrences may refer to the same effect identity. *Lineage* is the certified ancestry preserved when an edit carries or copies a workflow region. T, Δ, χ denote the active workflow, authorization log, and ordered record of resolved occurrences.

A pomset $p = (X_p, \prec_p)$ is a finite set of occurrences with a strict partial order. Its linearizations $\text{Lin}(p)$ are the words containing each member of X_p once and respecting $\prec_p$. An *indexed workflow contract* $\mathcal{C} = (O, \{p_o\}_{o \in O})$ is a finite nonempty family of pomsets. Each index o names one allowed completion outcome. Two outcomes remain distinct even if the calls still required for them later become identical. Write $X_{\mathcal{C}} = \bigcup_{o \in O} X_{p_o}$ and $\text{CL}(\mathcal{C}) = \bigcup_{o \in O} \text{Lin}(p_o)$. We require the runtime to distinguish completion from continuation.

$$v, w \in \text{CL}(\mathcal{C}) \wedge v \preceq_{\text{pref}} w \implies v = w.$$

This condition permits incomparable traces and different outcomes with the same trace, so it does not require deterministic outcomes. It requires a registered occurrence whenever continuing after a shared partial execution performs more work. The interface represents a choice between stopping and continuing by an authenticated terminal occurrence handled by the controller.

Contracts use the three partial constructors below. Each constructor requires occurrence-disjoint operands and is defined only when its result remains completion-observable.

$$\begin{aligned}
O_{C \oplus D} &= (\{L\} \times O_C) \uplus (\{R\} \times O_D), \\
p_{(L,o)} &= p_o, \quad p_{(R,v)} = p_v, \\
O_{C \otimes D} &= O_C \times O_D, \\
p_{(o,v)} &= p_o \uplus p_v, \\
O_{C \triangleright D} &= O_C \times O_D, \\
p_{(o,v)} &= p_o \triangleright p_v.
\end{aligned} \tag{2}$$

Before adding a registered or copied operand, the edit derivation deterministically renames its occurrences and extends ρ while retaining their authenticated effect identities. Well-formedness already makes carried operands disjoint. In equation (2), $\uplus$ adds no cross-order and $\triangleright$ orders every event of p before every event of q , while choice uses a tagged union of outcome indices and parallel and sequence use Cartesian products. Consequently, choice retains each arm's full indexed family, parallel pairs outcomes while permitting every causal linearization, and sequence adds a barrier. For raw prefix w , C/w retains each index whose pomset has a linearization prefixed by w , then removes w 's occurrences and incident order. It is undefined if no index remains. This operational "remaining contract" preserves outcome and occurrence identity.

Lemma 1 (Observable completion). *Every defined residual C/w is completion-observable, and either every retained pomset is empty or no retained pomset is empty.*

Proof. If a residual suffix v strictly prefixes v' , then wv strictly prefixes wv' in $\text{CL}(C)$. If one retained pomset is empty and another is nonempty, choose a nonempty linearization v of the latter. Then w strictly prefixes wv . Both contradict completion observability. $\square$

B. Authenticated execution record

1) *Recorded structure and active workflow:* The authenticated execution record is

$$H = (\omega, G, T, K, \Sigma_\zeta, \rho, \beta, \chi, \nu, \eta). \tag{3}$$

ω identifies the policy domain, and η identifies its current monitor version. G is a finite append-only record of workflow versions. Its node identities, base contracts, outcome-lineage authority rows, and typed edit records never change. Each continuation node b carries an immutable registered base contract Γ_b . Its branch cursor χ_b records which occurrences on that continuation have completed. The immutable records form a directed acyclic graph (DAG) whose edges record typed Fork, Restore, and Merge operations. T is the active workflow.

$$T ::= b:C \mid T \oplus_g T \mid T \parallel_g T \mid \text{jbar}(T, T) \mid T; T.$$

$s \in \{\text{open}, L, R\}$ records an unresolved choice or the arm that made recorded progress, and g groups exclusive or jointly active siblings. $\llbracket T \rrbracket$ uses both arms of an open choice and only the selected arm thereafter. It interprets parallel and jbar with

$\otimes$, and sequence with $\triangleright$. The unaddressed jbar marks a joined group and prevents another Merge.

T is complete when every pomset in $\llbracket T \rrbracket$ is empty. Lemma 1 excludes a residual containing both empty and nonempty pomsets. Set $\text{heads}(b:C) = \{b\}$. Open choice exposes both arms, selected choice its winner, parallel and jbar both arms, and sequence its left operand until completion and then its right. Completing an occurrence updates the work remaining at its leaf without deleting the recorded structure. Because sequence exposure tests complete, it retains completed left-outcome indices, so an isolated completed leaf remains addressable for later Fork or Restore. Joining removes g immediately but keeps its continuation behind a barrier until both arms complete. An occurrence is enabled when minimal in some pomset of $\llbracket T \rrbracket$, consistent with sequence exposure by observable completion. Occurrence disjointness and unique context decomposition give every enabled x a unique exposed leaf $\text{leaf}_T(x) = b$. For $a \in \{\text{new}(x, d), \text{reuse}(x, d)\}$, define the paired execution update

$$\text{HStep}((G, T), a) = (G \parallel (b, a), T/(x, a)), \quad b = \text{leaf}_T(x),$$

where $G \parallel (b, a)$ appends the progress record (b, a) , and $T/(x, a)$ residualizes the leaf and records any choice selection without storing a second cursor. Write $\text{HRes}((G, T), r)$ for iterating this update over a resolved word r .

$\text{addr}(T)$ follows the same recursion over branch and group identifiers. It includes exposed groups before descending into their active arms and only the exposed sequence operand, excluding inactive alternatives and unexposed operands.

K maps checkpoint identifiers to immutable earlier continuation records. Write $C_k = \text{contract}(K(k))$ for the residual contract stored by checkpoint k . Σ_ζ is a versioned immutable registry of typed edit schemas. ρ contains the authenticated occurrence-to-monitor-entry/effect-identity map, authorization tokens, lineage, scope, normalized tool requests, and request digests. For each current occurrence x , binding $\beta(x) = (j, \eta)$ names its monitor entry and version. χ is the ordered append-only record of resolved occurrences, including new and reuse bindings, and ν is the record version. Because χ records authorization and workflow progress rather than remote completion, a return value may still await settlement.

$\text{Save}(k, b)$ is an explicit checkpoint transition, not a seventh execution edit, and requires fresh k and addressable b . A trusted transaction stores an authenticated copy of continuation b 's remaining contract C_b , completed-occurrence cursor χ_b , and relevant call and outcome-authority rows. The checkpoint check requires $C_b = \Gamma_b/\text{raw}(\chi_b)$, and the transaction advances ν without changing the active workflow, authorization records, bindings, monitor versions, or permission trace. Both machines take this τ -transition, which makes a concurrent installation stale because installation authenticates H . Restore accepts only records created by this rule.

2) Well-formedness:

Definition 1 (Well-formed execution record). *H is well formed when the following groups hold.*

- 1) Structure. G is acyclic, and continuation and group identifiers are globally unique across the recorded DAG and current workflow. Every addressable identifier has a unique context decomposition, and every head of T is a current maximal node. Every checkpoint names an immutable node coordinate and a value-snapshotted residual/cursor pair derived from G . Every root or schema-introduced outcome lineage has exactly one signed policy-authority row, preserved by context identity, carry, clone, and checkpoint snapshots.
- 2) Identity and progress. Occurrence and monitor-entry identifiers are unique, while occurrences mapped to the same effect identity agree on the call authorization service, scope, normalized request, digest, and lineage. Resolved occurrences are downward closed in their pomset, consistent with every choice status, and absent from the remaining workflow.
- 3) Current workflow. $\text{dom}(\beta) = X_{[T]}$, and every current occurrence x has exactly one authenticated binding $\beta(x) = (j_x, \eta)$ to the domain's single current monitor version. Every current leaf $b:C$ is exactly the indexed residual of its immutable registered base Γ_b by the raw projection of its resolved cursor χ_b , rather than a caller-pruned subfamily. Every constructor in T is defined, so every subtree interpretation is completion-observable.
- 4) Schema and domain. Every edge and introduced contract is justified by one immutable schema at version ζ . Every current, checkpoint, or schema occurrence, effect identity, monitor entry, lineage, and scope belongs to ω . Effect identities, authorization logs, and policies of distinct domains are disjoint, and an edit never crosses a domain.

The trusted installation service additionally requires that the policy domain's one current monitor version is active.

3) *Authorization log and security core:* Outside H , the model stores policy version κ , a regular prefix-closed permission policy $\mathcal{A} \subseteq \mathcal{U}^*$ over label alphabet $\mathcal{U}$, append-only authorization log Δ , and ordered dispatch queue out . Authorization records can only be appended. Settlement advances an authorization record's status and fills its return record with a stable return value. Each authorization record binds its effect identity, immutable normalized tool request, creator occurrence, and permission label. Each entry in χ points to the authorization record created by its new event or, for reuse, to an earlier authorization record, making retry and reuse return lookup functions of (χ, Δ) . $\text{lab}(\Delta)$ is the permission trace: the ordered labels of calls that create authorization records. Restore may reconstruct workspace payload but never restores Δ .

A security core $(\kappa, H, \Delta, \text{out}, \mathcal{A})$ is *well formed* when H is well formed, $\text{lab}(\Delta) \in \mathcal{A}$, each effect identity has at most one authorization record created by new, each new entry names that record, each reuse entry names an earlier authorization record, the authorization log, dispatch queue, and policy are scoped to ω , and the queue preserves new order. This joint invariant is assumed below.

C. Adversary and trusted boundary

The adversary controls Agent output, edit requests, tool arguments, copied workspace, retries, stale authorization tokens, and scheduling, including races between calls and edits. It cannot forge registries, the checkpoint directed acyclic graph (DAG), effect-identity certificates, policy, or authorization log. The policy authority signs policies, outcome requirements, and authorized removals; the call authorization service registers protected calls, effect identities, and authorization tokens. The trusted controller derives and checks each edit, while the trusted installation service re-verifies the answer and atomically replaces the monitor. These roles, the authenticated registries and authorization log, certificate verifier, atomic rules, and complete mediation form the trusted computing base (TCB). Within one policy domain, this TCB completely mediates every protected call and provides a linearizable, crash-stable transaction for every modeled state change and visible edit answer. A recovered labeled transition system (LTS) state shows each transaction entirely uncommitted or committed according to whether its linearization point occurred. Recovered states expose every transaction as entirely uncommitted or committed, while compiler caches are volatile. Before the checker takes a source snapshot, the call authorization service may instantiate finitely many pre-authorized schema templates. For tool request m , it computes normalized request $m^\circ = \text{canon}_\kappa(m)$, reuses only an effect identity certified for m° or allocates a fresh one, and signs $(x, j, d, \text{scope}, \text{digest}(m^\circ), \text{lineage})$ into the authenticated registry. The resulting finite registered call model R remains unchanged while that snapshot is checked. Unmatched calls fail closed. The platform supplies authenticated $\mathcal{W}, R$, edit schemas, normalization rules, policy, checkpoints, and one consistent snapshot of the authorization log and dispatch queue. From these inputs and an identifier-only Agent request, the controller derives the edited workflow, required outcomes, resolved calls, fixed point, monitor, certificate, monitor version, and authorization tokens. The Agent may supply or weaken none of these derived objects.

A boundary-finite typed Agent workflow has only finitely many distinct protected-call traces that reach an outcome. CheckReg accepts a finite registered call model only when it faithfully represents the outcomes, call order, effect identities, and first-time authorization records of that workflow. The model stores these traces as outcome-indexed runs $\text{BRuns}(R)$ and a map λ from source traces; write $\mathcal{W} \sqsubseteq_\lambda R$ when λ is a bijection preserving outcomes, the mapping from call positions to effect identities, order, outcome authority, and authorization-log updates. Definition 6 and lemma 13 give the exact finite check and proof. Registered contracts enter through CheckReg, and remote exactly-once effects enter through authenticated receipts from an idempotent or queryable sink.

D. Security Objective

The security objective is fixed independently of the checker that will implement it. For an edit derived from an authenticated source snapshot, a safe realization must satisfy four requirements. First, it preserves every source outcome that is

neither already satisfied nor explicitly authorized for removal, and every admitted execution belongs to the derived target workflow. Second, every admitted prefix is allowed by the policy after accounting for authorization records already present in the append-only log. Third, each admitted prefix extends to a complete admitted execution. Fourth, every outcome still compatible with an admitted prefix retains such a completion. The last requirement rules out the prefix-unsafe example in section II, even though each outcome has a safe trace in isolation. Definition 2 states these requirements as an indexed language relation; lemma 4 then shows that the checker computes its greatest member rather than defining safety by its own output.

IV. EXACT CHECKING OF EXECUTION EDITS

The checker first derives the only admissible target workflow from the authenticated source and the requested edit. It then resolves call occurrences against existing authorization records and removes precisely those complete executions that violate the security objective after some prefix.

A. Edit Derivation and Outcome Preservation

1) *Edit certificates*: The six disjoint constructors of Edit_6 are exactly the typed operation signatures in table I. The set excludes Save, Root, Extend, protected calls, and caller-supplied targets.

An untrusted request u supplies an edit kind, current object identifiers, and schema identifier σ , never a target contract, while Σ_ζ binds σ to that kind, exact source and checkpoint fingerprints, introduced suffix contracts, a lineage certificate, and the signed authority assignment for each introduced lineage.

The edit certificate defined below makes the derivation checkable. For indexed source $\mathcal{C}$, schema-designated carried or cloned target region $\mathcal{D}$, and parent map $\mu : X_{\mathcal{D}} \rightarrow X_{\mathcal{C}}$, write $\mathcal{C} \preceq_\mu \mathcal{D}$ exactly when μ is a global bijection, some surjection $\pi : O_{\mathcal{D}} \twoheadrightarrow O_{\mathcal{C}}$ satisfies the conditions below for every target outcome v , and μ preserves effect identity and lineage.

$$\begin{aligned} x \in X_{p_v} &\iff \mu(x) \in X_{p_{\pi(v)}}, \\ x \prec_v y &\iff \mu(x) \prec_{\pi(v)} \mu(y). \end{aligned}$$

Global bijection and occurrence–outcome membership equivalence prevent splitting a source occurrence shared across outcomes into outcome-specific target occurrences.

For candidate T' , let R_u index the disjoint carried or cloned target regions inside the replaced subtree. An edit certificate $\mathcal{P}_u = \{(\mathcal{C}_r, \mathcal{D}_r, \mu_r, \pi_r)\}_{r \in R_u}$ contains the relation for each indexed region. A separate identity map checks that the exposed one-hole context preserves syntax, outcome indices, occurrences, effect identities, lineage, and order. Writing X_{ctx} for its occurrences and $X_{\text{suf}(u)}$ for schema-declared suffix occurrences, the checker verifies the following disjoint cover.

$$X_{[T']} = X_{\text{ctx}} \uplus \biguplus_{r \in R_u} X_{\mathcal{D}_r} \uplus X_{\text{suf}(u)}.$$

The row constructor and exposed-context syntax induce $\text{pr}_r^u : O_{[T']} \rightarrow O_{\mathcal{D}_r}$. For choice, this projection is defined only in

the arm containing r , while product and sequence select the corresponding indexed component. Define

$$\begin{aligned} \text{RegionReq}_u &= \{(r, o) \mid r \in R_u, o \in O_{\mathcal{C}_r}\}, \\ \text{Lift}_u(t, r, o) &\iff \exists v \in O_{\mathcal{D}_r}. \text{pr}_r^u(t) = v \wedge \pi_r(v) = o. \end{aligned}$$

Let O_{byp} be source-context outcomes whose alternatives bypass the hole, and let $\iota_{\text{ctx}}^u : O_{\text{byp}} \hookrightarrow O_{[T']}$ be the injection induced by context identity. Define the single required-outcome set and its bypass lift by

$$\begin{aligned} \text{Req}_u &= \text{RegionReq}_u \uplus (\{\text{ctx}\} \times O_{\text{byp}}), \\ \text{Lift}_u(t, (\text{ctx}, o)) &\iff t = \iota_{\text{ctx}}^u(o). \end{aligned} \quad (4)$$

For regional $a = (r, o) \in \text{RegionReq}_u$, $\text{Lift}_u(t, a)$ denotes the relation above. The checker also requires $\forall a \in \text{Req}_u. \exists t \in O_{[T']}$. $\text{Lift}_u(t, a)$, ensuring a full-target lift for every required source outcome, no outcome-specific splitting of shared occurrences, and no inherited occurrence mislabeled as a suffix.

The schema-verification phase checks the following judgment.

$$\Sigma_\zeta; H \vdash_{\text{edit}} u \Rightarrow (T', \mathcal{P}_u)$$

It holds exactly when every named current object belongs to $\text{addr}(T)$, every introduced object remains scoped to ω , the source fingerprint matches, the applicable $\mathcal{P}_u$, disjoint cover, and Lift_u requirements in table I hold, any removal is authorized, and $[T']$ is defined. Authorized removal is confined to the unselected arm of a typed MergeSelect or the current continuation of RestoreReplace, using the exact authority-bound sets below rather than Agent declarations.

2) *Required-outcome preservation*: The source row itself, rather than a caller-proposed target, determines every outcome that the edit must account for. Fork accounts for both copies, Restore for the current and checkpoint continuations, MergeSelect for the selected and unselected arms, and MergeJoin for both arms. Outcomes in the surrounding context that bypass the edited region are unchanged. Definition 4 defines these source outcomes, completed outcomes, and the only outcomes eligible for signed removal. The checker accepts the schema only when the target-derived set from equation (4) satisfies

$$\begin{aligned} \text{Req}_u &= \text{SourceReq}(H, u) \setminus (\text{Completed}_{H,u} \uplus \text{Removed}_{H,u}), \\ \text{and hence the exact disjoint partition} \end{aligned}$$

$$\text{SourceReq}(H, u) = \text{Req}_u \uplus \text{Completed}_{H,u} \uplus \text{Removed}_{H,u}.$$

Thus the Agent cannot make an edit pass by quietly dropping an inconvenient outcome. Checked lifts preserve effect identity, lineage, and order, while installation retains recorded evidence for every already-satisfied source outcome and the exact signed removal set.

Let $\mathcal{E}[-]$ be a one-hole *exposed* active-workflow context, so its hole has the addressability property above. carry preserves continuation, outcome, and occurrence identifiers, pomset order, effect identities, lineage authority, base contracts, current branch cursors, and nested choice status. clone_r gives continuation, outcome, and occurrence identifiers deterministic fresh names while preserving pomset order and authenticated effect identity,

scope, digest, lineage, and lineage authority. It registers each copied remainder as a new base contract with an empty cursor. Schema suffix leaves likewise start with their registered base contract and an empty cursor. The schema-derived contracts $E_\sigma^L, E_\sigma^R, E_\sigma^J$ are suffixes, not replacements. Define

$$L_\sigma = \text{clone}_L(\mathcal{C}) \triangleright E_\sigma^L, \quad R_\sigma = \text{clone}_R(\mathcal{C}) \triangleright E_\sigma^R.$$

New identifiers and the target monitor version η^+ are deterministically allocated from $(\omega, \zeta, \nu, u, \text{role})$. role identifies the edit row and operand position. Write $\text{Cl}_r(\mathcal{C})$ for the certified tuple $(\mathcal{C}, \text{clone}_r(\mathcal{C}), \mu_r, \pi_r)$, and $\text{Ca}(\mathcal{C})$ for $(\mathcal{C}, \text{carry}(\mathcal{C}), \mu, \pi)$. Each abbreviation asserts the corresponding $\preceq_\mu$ relation. For an active workflow T_0 , $\text{Ca}(T_0)$ abbreviates $\text{Ca}(\llbracket T_0 \rrbracket)$ and has target contract $\llbracket \text{carry}(T_0) \rrbracket$. Every b or g named below must lie in $\text{addr}(T)$.

In an open choice, the first new or reuse atomically selects its arm and disables later occurrences from the other, so $\text{MergeSelect}(g, w, \sigma)$ has no derivation after any recorded progress on the unselected arm, including progress that creates no authorization record. MergeJoin replaces the addressable $\llbracket g$ node with jbar , removes g from the addressable workflow, and prevents the same pair from being joined again.

The model replaces the monitor for the whole policy domain rather than disabling only selected monitor entries, so every current entry in H belongs to $\eta^- = \eta$.

3) *Monitor-version replacement*: The execution-edit judgment is the least relation generated by one row of table I inside a well-typed single-hole context $\mathcal{E}$, followed by the common update below; no other execution-edit derivation exists. Let e_u contain its typed edge and exactly the newly allocated clone, suffix, and group nodes prescribed by T' ; carried node identities remain unchanged. For every current target occurrence x , the trusted call authorization service deterministically mints j_x^+, h_x^+ , and ρ' preserves the certified effect identity, lineage, scope, digest, and records needed for Retry while refreshing monitor entries and authorization tokens. Define

$$\mathcal{H}_u^+ = \{(x, j_x^+, h_x^+) \mid x \in X_{\llbracket T' \rrbracket}\}, \\ \beta'(x) = (j_x^+, \eta^+).$$

The new authorization tokens are inactive and inaccessible to the caller before installation. The complete judgment is shown below.

$$\begin{aligned} \Sigma_\zeta; H \vdash u \Downarrow (H', \mathcal{C}_{H,u}, \eta^-, \eta^+), \\ H' = (\omega, G \cdot e_u, T', K, \Sigma_\zeta, \rho', \beta', \chi, \nu + 1, \eta^+), \\ \mathcal{C}_{H,u} = \llbracket T' \rrbracket, \quad \eta^- = \eta. \end{aligned} \quad (5)$$

Here $G \cdot e_u$ retains source progress, authority rows, and old nodes, and initializes each clone or suffix with its residual base, inherited or schema-signed authority rows, and no progress. Thus an edit preserves K, Σ_ζ, χ , appends one typed provenance edge and its target nodes, refreshes every current monitor entry and authorization token, and moves the domain to η^+ . Table constructors determine new group modes, while carried nested modes remain unchanged. A request naming a nonexistent checkpoint, wrong group mode, nonmember winner,

mismatched schema fingerprint, unauthorized removal, or undefined composition has no derivation. Independent domains step by asynchronous product, and no theorem allows disabling only a selected subset of one domain's monitor entries.

Lemma 2 (Schema fidelity and complete monitor replacement). *For well-formed H , the judgment in equation (5) is deterministic. If it derives H' , then H' is well formed. It preserves required outcomes and their authority rows: evidence for already-satisfied outcomes remains recorded, and no pending source lineage disappears outside $\text{Removed}_{H,u}$. The identity-preserved context and family $\mathcal{P}_u$ preserve global occurrence sharing, effect identity, lineage, lineage authority, and causal order. For every $a \in \text{Req}_u$, some full target outcome t satisfies $\text{Lift}_u(t, a)$, including the identity lift of every bypassing context outcome. Every current target occurrence has a monitor entry bound to η^+ and one authorization token in $\mathcal{H}_u^+$, and none remains bound to η^- .*

Proof sketch. The six row cases use certified clone/carry bijections and context identity to preserve sharing, effect identity, lineage, and order and to construct every Lift_u witness. The common poststate establishes well-formedness and binds every current target occurrence to η^+ ; section A-D1 gives the full argument. $\square$

Recall the permission-label alphabet $\mathcal{U}$ and regular prefix-closed policy $\mathcal{A} \subseteq \mathcal{U}^*$. For the authenticated authorization log Δ , $\text{lab}(\Delta) \in \mathcal{A}$ is its permission trace and D_Δ is the set of effect identities that already have authorization records. Fix $\Theta = (\kappa, H, \Delta, \text{out}, \mathcal{A}, u)$ and write $\text{src}(\Theta) = (\kappa, H, \Delta, \text{out}, \mathcal{A})$. When defined, the tagged root, execution-edit, or registered-extension derivation is unique. Write $\text{Derive}(\Theta) = (\kappa_u, H_u, \mathcal{A}_u, \mathcal{C}_u, \eta^-, \eta^+)$. Root derivation additionally requires that $\text{CheckReg}(\mathcal{W}, R)$ accept. An execution edit has $(\kappa_u, \mathcal{A}_u) = (\kappa, \mathcal{A})$ and uses the unique execution-edit judgment. $\text{Extend}(\xi)$ uses the authenticated judgment in section III. Let ρ_u be the authenticated target registry in H_u , let $O_{H,u} = O_{\mathcal{C}_u}$, and use $H' = H_u$ and $\rho' = \rho_u$ where earlier notation is convenient. All languages below depend on Θ , the schema registry, ρ_u , and target policy $\mathcal{A}_u$, but not on the dispatch queue out . The executable checker nevertheless binds out into the authenticated source snapshot used for installation. We write only H, u where the remaining parameters are fixed.

B. new and reuse resolution

Resolution never erases an occurrence retained by the derived contract. Starting with $D_0 = D_\Delta$, the deterministic transducer $\text{Resolve}_{\Delta}^{\rho_u}$ maps a raw word $w = x_1 \cdots x_n$ to a resolved word $z_1 \cdots z_n$. For $d_i = q_{\text{eff}}^{\rho_u}(x_i)$,

$$z_i = \begin{cases} \text{new}(x_i, d_i), & d_i \notin D_{i-1}, \quad D_i = D_{i-1} \cup \{d_i\}; \\ \text{reuse}(x_i, d_i), & d_i \in D_{i-1}, \quad D_i = D_{i-1}. \end{cases} \quad (6)$$

The permission projection erases reuse symbols and maps each new symbol to its effect identity's permission label.

$$\text{auth}(\text{new}(x, d)) = \ell_{\rho_u}(d), \quad \text{auth}(\text{reuse}(x, d)) = \epsilon.$$

TABLE I
SCHEMA-CHECKED AGENT EXECUTION EDITS. $E_\sigma^\bullet$ IS DERIVED FROM Σ_ζ , NEVER SUPPLIED AS A TARGET CONTRACT.

Operation u	Required current shape	Derived target shape	Checked preservation
ForkChoice(b, σ)	$\mathcal{E}[b:C]$	$\mathcal{E}[b_L:L_\sigma \oplus_g^{\text{open}} b_R:R_\sigma]$	$\{\text{Cl}_L(C), \text{Cl}_R(C)\}$
ForkParallel(b, σ)	$\mathcal{E}[b:C]$	$\mathcal{E}[b_L:L_\sigma \parallel_g b_R:R_\sigma]$	$\{\text{Cl}_L(C), \text{Cl}_R(C)\}$
RestoreReplace(b, k, σ)	$\mathcal{E}[b:C], k \in \text{dom } K$	$\mathcal{E}[b_k:\text{clone}_k(C_k)]$	$\{\text{Cl}_k(C_k)\}$; current continuation removed
RestoreLive(b, k, σ)	$\mathcal{E}[b:C], k \in \text{dom } K$	$\mathcal{E}[b:\text{carry}(C) \parallel_g b_k:\text{clone}_k(C_k)]$	$\{\text{Ca}(C), \text{Cl}_k(C_k)\}$
MergeSelect(g, w, σ)	$\mathcal{E}[T_L \oplus_g^* T_R], s \in \{\text{open}, w\}$	$\mathcal{E}[\text{carry}(T_w); b_J:E_\sigma^J]$	$\{\text{Ca}(T_w)\}$; other arm removed
MergeJoin(g, σ)	$\mathcal{E}[T_L \parallel_g T_R]$	$\mathcal{E}[\text{bar}(\text{carry}(T_L), \text{carry}(T_R)); b_J:E_\sigma^J]$	$\{\text{Ca}(T_L), \text{Ca}(T_R)\}$

For resolved z , $\text{raw}(z)$ drops the new/reuse mode and effect identity but preserves the ordered occurrence identifiers. A Retry repeats the same normalized request and adds no occurrence. Restored occurrences sharing an effect identity remain distinct and may resolve differently.

Lemma 3 (Resolution). *Resolve $_{\Delta}^{\rho_u}$ is deterministic, length preserving, and prefix compatible. Its new projection contains each effect identity at most once outside D_Δ , while its occurrence projection is exactly the input word. Moreover, if $\text{Resolve}_{\Delta}^{\rho_u}(rv) = zv'$, then $z = \text{Resolve}_{\Delta}^{\rho_u}(r)$ and $v' = \text{Resolve}_{\Delta_z}^{\rho_u}(v)$, where Δ_z extends Δ by the new events of z .*

Proof. Induction on input length shows that the transducer emits one prefix-determined output symbol per input symbol and updates D_i only for new. Set membership ensures one new emission per effect identity. Both cases retain x_i . Splitting after r gives the streaming equation. $\square$

Thus a retry changes neither the resolved call record nor the permission trace, although its successful API reply is observable as $\text{Return}(t, [v]_{\cong})$. A copied occurrence that resolves as reuse advances the call record without changing the permission trace.

C. Safe Completion After Every Prefix

1) *Largest family safe after every prefix:* For each derived outcome $o \in O_{H,u}$, define all of its completed executions and the subset allowed by policy.

$$\begin{aligned} R_o(H, u) &= \{\text{Resolve}_{\Delta}^{\rho_u}(w) \mid w \in \text{Lin}(p_o)\}, \\ M_o(H, u) &= \{z \in R_o(H, u) \mid \text{lab}(\Delta)\text{auth}(z) \in \mathcal{A}_u\}. \end{aligned} \quad (7)$$

For a resolved prefix z , let

$$\text{Compat}(z) = \{o \mid \text{raw}(z) \preceq_{\text{pref}} w \text{ for some } w \in \text{Lin}(p_o)\}.$$

Relation Preserve requires every source outcome that is neither completed nor authorized for removal to have a checked lift into $\hat{B}$; definition 5 gives its full definition.

Definition 2 (Indexed realization). *A generated language and indexed marked family $(L, \hat{B})$ realize (H, u) when all conditions below hold.*

- 1) $\hat{B} \neq \emptyset$;
- 2) source and target fidelity: $\text{Preserve}(H, u, \hat{B})$,
 $\hat{B} \subseteq \biguplus_{o \in O_{H,u}} \{o\} \times R_o(H, u)$, and $L \subseteq \text{Pref}(\bigcup_{o \in O_{H,u}} R_o(H, u))$;
- 3) policy safety: $\text{lab}(\Delta)\text{auth}(z) \in \mathcal{A}_u$ for every $z \in L$;

- 4) completion nonblocking: $L = \text{Pref}(\pi_2 \hat{B})$; and
- 5) persistent outcome coverage: for every $z \in L$ and $o \in \text{Compat}(z)$, some $(o, b) \in \hat{B}$ extends z .

Checking each outcome separately is insufficient. After a partial execution shared by several outcomes, every outcome still compatible with that partial execution must retain a policy-safe completion. The operator below removes exactly the completed executions that violate this condition. Put

$$\hat{B}_0 = \biguplus_{o \in O_{H,u}} \{o\} \times M_o.$$

For $\hat{B} \subseteq \hat{B}_0$, define the monotone operator

$$\hat{\Phi}(\hat{B}) = \left\{ (o, b) \in \hat{B}_0 \mid \begin{array}{l} \forall z \preceq_{\text{pref}} b. \forall v \in \text{Compat}(z). \\ \exists (v, b') \in \hat{B}. z \preceq_{\text{pref}} b' \end{array} \right\}. \quad (8)$$

Set $\hat{B}_{i+1} = \hat{\Phi}(\hat{B}_i)$. Since $\hat{B}_1 \subseteq \hat{B}_0$, monotonicity yields a finite descending chain, whose limit is

$$\begin{aligned} \hat{B}_{H,u}^\dagger &= \nu \hat{B}. \hat{\Phi}(\hat{B}) = \bigcap_{i \geq 0} \hat{B}_i, \\ B_{H,u}^\dagger &= \pi_2 \hat{B}_{H,u}^\dagger, & W_{H,u}^\dagger &= \text{Pref}(B_{H,u}^\dagger). \end{aligned}$$

Definition 3 (Safe execution edit). *If $\text{Derive}(\Theta)$ is undefined, then $\text{SafeEdit}(\Theta)$ is false. Otherwise,*

$$\text{SafeEdit}(\Theta) \iff \hat{B}_{H,u}^\dagger \neq \emptyset. \quad (9)$$

$\text{SafeEdit}(H, u)$ fixes Θ and both registries. At $z = \epsilon$, nonemptiness covers every required outcome. The existential quantifier chooses a call order without assuming fairness, and the universal condition keeps every compatible outcome completable after every prefix.

Lemma 4 (Largest safe realization). *A realization exists iff $\text{SafeEdit}(H, u)$. When the edit is safe, $(W_{H,u}^\dagger, \hat{B}_{H,u}^\dagger)$ is a realization, and every realization $(L, \hat{B})$ satisfies $\hat{B} \subseteq \hat{B}_{H,u}^\dagger$ and $L \subseteq W_{H,u}^\dagger$.*

Proof. Fidelity, completion nonblocking, and safety place any realization's indexed marked family inside $\hat{B}_0$. Persistent coverage makes it post-fixed, so monotonicity and finite Knaster-Tarski place it inside $\hat{B}^\dagger$, and projected prefixes give $L \subseteq W^\dagger$. Conversely, $\hat{B}^\dagger = \hat{\Phi}(\hat{B}^\dagger)$ supplies persistent coverage. Together with checked lifts, this coverage establishes Preserve at ϵ . The inclusion $\hat{B}^\dagger \subseteq \hat{B}_0$ supplies target fidelity and safe marked

words, while prefix closure of $\mathcal{A}_u$ supplies prefix safety. Finally, $W^\dagger = \text{Pref}(\pi_2 \hat{B}^\dagger)$ supplies nonblocking. Hence the fixed point is a realization exactly when nonempty. $\square$

$\text{GReal}(\Theta; L, \hat{B})$ denotes a realization containing every other one componentwise, and its unique witness is $\text{Spec}(\Theta)$. Lemma 4 proves that this witness exists exactly when $\hat{B}_{H,u}^\dagger \neq \emptyset$ and then equals $(W_{H,u}^\dagger, \hat{B}_{H,u}^\dagger)$.

2) *Outcome-Wise Checking Counterexample*: Section II introduced the contract $X \otimes (Y \oplus Z)$ under policy $\text{Pref}\{yx, xz\}$. For its resolved calls, write $\bar{x} = \text{new}(x_o, d_x)$, $\bar{y} = \text{new}(y_o, d_y)$, and $\bar{z} = \text{new}(z_o, d_z)$. The candidate family initially contains $\bar{y}\bar{x}$ and $\bar{x}\bar{z}$. The first fixed-point iteration removes $\bar{x}\bar{z}$. The second iteration removes $\bar{y}\bar{x}$, because after the first removal the empty prefix no longer has a completion for the $X \otimes Z$ outcome. The empty fixed point is the formal reason the exact checker rejects.

Proposition 1 (Outcome-indexed nonblocking correspondence). *For nonempty $\hat{B} \subseteq \hat{B}_0$, let $G_{\hat{B}}$ be the prefix trie of $L_{\hat{B}} = \text{Pref}(\pi_2 \hat{B})$. Mark state z by α_o exactly when $o \in \text{Compat}(z)$, and mark state b by ω_o exactly when $(o, b) \in \hat{B}$. Then*

$$\hat{B} \subseteq \hat{\Phi}(\hat{B}) \\ \iff \forall o \in O_{H,u}. G_{\hat{B}} \text{ is } (\alpha_o, \omega_o)\text{-nonblocking}.$$

Consequently, if $\hat{B}^\dagger$ is nonempty, it is the largest nonempty completion subfamily of $\hat{B}_0$ satisfying these outcome-specific nonblocking conditions. If it is empty, no nonempty subfamily satisfies the conditions. Ordinary marker nonblocking of the unaugmented projected language is strictly weaker.

Proof. From reachable prefix z , an ω_o -marked state is reachable exactly when some $(o, b) \in \hat{B}$ extends z . Quantifying over exactly the states marked α_o is therefore the post-fixed-point condition. Greatestness follows from the greatest-fixed-point construction. For strictness, the preceding instance has the marker-nonblocking projected pair $(\text{Pref}\{\bar{y}\bar{x}, \bar{x}\bar{z}\}, \{\bar{y}\bar{x}, \bar{x}\bar{z}\})$, but its indexed fixed point is empty. $\square$

Proposition 1 gives the exact interface to generalized nonblocking [15], [16]. After the trusted controller derives α_o, ω_o from the requested edit, authorized removals, effect identities, authorization records, and source snapshot, a generalized-nonblocking solver computes the same largest family. Complete mediation makes new/reuse controllable. Uncontrollable resolved events compose this construction with the standard supremal-controllable refinement. Our theory additionally constructs the most permissive Agent monitor or an impossibility certificate, carries the result through atomic installation, and proves that neither a caller-proposed workflow nor its traces contain enough information to reproduce the derivation.

3) *Rejection Certificates*: Each removed (o_{rem}, b) records its rank i , a prefix $z \preceq b$, and uncovered $o_{\text{miss}} \in \text{Compat}(z)$. Induction shows that every post-fixed P is contained in every $\hat{B}_i$. An empty final family therefore excludes every realization,

and installation recomputes the ranked rejection certificate and the remaining witnesses.

4) *Residual Coherence*: To show that the checker can be applied again after each permitted partial execution, fix $H, u, \mathcal{A}_u$ after deriving $\mathcal{C}$, and write $W^\dagger(\mathcal{C}, \Delta)$ and $\hat{B}^\dagger(\mathcal{C}, \Delta)$ with those parameters suppressed.

Lemma 5 (Residual coherence). *If $z \in W^\dagger(\mathcal{C}, \Delta)$, then $\text{Resolve}_\Delta^{\rho'}(\text{raw}(z)) = z$, $\text{lab}(\Delta_z) = \text{lab}(\Delta)\text{auth}(z)$, and*

$$W^\dagger(\mathcal{C}, \Delta)/z = W^\dagger(\mathcal{C}/\text{raw}(z), \Delta_z), \\ \hat{B}^\dagger(\mathcal{C}, \Delta)/z = \hat{B}^\dagger(\mathcal{C}/\text{raw}(z), \Delta_z),$$

where $\hat{B}/z = \{(o, v) \mid (o, zv) \in \hat{B}\}$.

Proof sketch. Streaming resolution gives raw recoverability, the authorization-log equality, and equality of the source and residual candidate families. Residualizing $\hat{B}^\dagger$ gives a post-fixed family for the successor, while prefixing any successor post-fixed family by z gives one for the source; greatestness yields the fixed-point equality, and $\text{Pref}(B)/z = \text{Pref}(B/z)$ yields the generated-language equality. The full proof is in section A-D3. $\square$

Thus every accepted prefix retains a safe completion for every residual outcome index still contract-compatible with that prefix.

5) *Checker Outputs*: The checker returns Invalid when derivation fails, Reject with a ranked impossibility certificate when the fixed point is empty, and Installable with the greatest realization otherwise. Each answer carries a fingerprint of the complete authenticated state that was checked. Section A-C gives the exact payloads and verification rules. Proposition 2 proves an output-sensitive $O(D+n\Lambda+|\mathcal{G}_{\text{cov}}|+V_B \log(2+V_B))$ -time, $O(D+\Lambda+|\mathcal{G}_{\text{cov}}|)$ -space bound. Exponential growth occurs only in explicitly emitted linearizations and prefix-required-outcome support pairs.

V. RUNTIME ENFORCEMENT

The safety checker produces either an impossibility certificate or the largest safe family of completed executions. The trusted installation service turns a nonempty family into runtime enforcement in three steps. It builds a monitor for the family's partial executions, binds that monitor to the authenticated source snapshot, and atomically installs it under a new monitor version. The protocol serializes every overlapping protected call either before or after installation. The proof compares this implementation with an ideal machine that installs an accepted specification in one atomic step and admits exactly its resolved prefixes. The ideal transition system and the field-level implementation rules appear in section A-H and table III; this section gives the monitor construction, the installed-state invariant, and the installation boundary needed by the main theorem.

A. Canonical Occurrence-Sensitive Monitor

The greatest realization has the finite, canonically installable projected pair $(W, B) = (W_{H,u}^\dagger, B_{H,u}^\dagger)$. For $L/z = \{v \mid zv \in L\}$, use residual states $q_z = (W^\dagger/z, B^\dagger/z)$.

$$\begin{aligned} Q_{W,B} &= \{q_z \mid z \in W\}, \\ q^0 &= q_\epsilon = (W^\dagger, B^\dagger), \\ q_z &\xrightarrow{a} q_{za} \iff za \in W. \end{aligned} \quad (10)$$

Call this monitor automaton $\text{Can}(W, B)$. Every state is reachable and is marked exactly when $\epsilon \in B^\dagger/z$. Each transition symbol a contains occurrence x , effect identity d , and its new or reuse mode, so one call position cannot use another position's monitor entry.

Lemma 6 (Exact compiled realization). *The generated and marked languages of the program in equation (10) are exactly $(W_{H,u}^\dagger, B_{H,u}^\dagger)$. The program is the smallest deterministic marked partial automaton for this pair up to state renaming.*

Proof. Induction on z shows that the automaton reaches $(W^\dagger/z, B^\dagger/z)$, enables a iff $za \in W^\dagger$, and marks the reached state iff $z \in B^\dagger$. Distinct states differ on a generated or marked suffix and must remain distinct. The automaton is the paired-language derivative construction [19]. $\square$

B. Atomic Installation

At installation, the service re-derives the edit and recomputes the fixed point from the current authenticated state. If any protected call has committed since compilation, the candidate is stale and installs nothing. Otherwise, one policy-domain transaction closes the old monitor version, activates the canonical monitor, replaces every current occurrence binding and authorization token, and only then exposes the successful edit answer. Protected calls verify their occurrence, effect identity, normalized request, authorization token, and current monitor transition in the same transaction that records progress. A new call creates one authorization record and queues the normalized request; an reuse call advances the workflow by linking to the existing record without extending the permission trace.

a) Installed-state invariant: We write $\text{AgentSec}(S; c, r)$ when five conditions hold: the installed workflow and its required outcomes come from the checked edit; the stored nonempty execution family is the checker's largest safe family; the workflow, authorization log, queue, and checkpoints match the executed prefix r ; the current monitor is exactly the derivative after r ; and exactly one monitor version is active, with one current entry and authorization token for every remaining call position. For a submitted tool request, one transaction verifies the call position, effect identity, normalized request, authorization token, current monitor version, and enabled monitor transition before recording any progress. The same transaction either creates the authorization record and queues the request, or records reuse of an existing authorization and return value.

b) Installation check: Write $\text{Ready}(S, u, Y)$ when the service re-derives u from the current state, recomputes the same nonempty fixed point and canonical monitor carried by Y , matches Y 's checked-state fingerprint, confirms the old monitor version is current, and confirms the new version is not active. If these checks pass, one transaction closes the old version, installs and activates the new monitor, replaces every current entry and authorization token, and then exposes the successful edit answer. A changed state makes the candidate stale and forces a new check. The complete field updates, retries, dispatch, settlement, and recovery rules appear in section A-I and table III.

Lemma 7 (Edit-call serialization). *In every reachable well-formed compiled state, a visible edit answer and a protected call in the same policy domain have a unique order at the installation point. If the call commits first, it changes the checked state, so the earlier fingerprint cannot install. If Install commits first, the old call cannot resolve new or reuse. Steps in independent domains commute by product semantics.*

Proof. If the call commits first, its recorded progress changes the checked-state fingerprint. If installation commits first, it closes the old monitor version, so the old authorization token fails. Atomicity excludes a partial third outcome; independent domains commute. $\square$

C. Trust Boundary

The TCB is the trusted boundary defined in section III. Search, minimization, and preloading remain untrusted because installation checks them again. The next section states the exactness and runtime guarantees obtained at this boundary.

VI. FORMAL GUARANTEES

The theorem connects authenticated edit derivation, exact checking, monitor construction, atomic installation, and every later runtime step. Its central clause says that the checker accepts if and only if a safe implementation exists; the other clauses establish faithful registration, necessary state information, repeated use, and runtime correctness.

Fix $\delta = (\kappa_u, H_u, \mathcal{A}_u, C_u, \eta^-, \eta^+)$. An *installed realization* is a pair $(M, \bar{B}_M)$ whose generated and marked languages realize (H, u) as in definition 2, installed atomically at η^+ while η^- is closed and all target monitor entries and tokens are replaced.

a) Information required end to end: At a checked source snapshot, $\text{Sec}(S; c, r) = \mathbf{P} \bowtie \mathbf{I} \bowtie \mathbf{E} \bowtie \mathbf{C}$ joins workflow structure and schemas ($\mathbf{P}$), call-effect identity ($\mathbf{I}$), effect progress ($\mathbf{E}$), and monitor-version and installation coordinates ($\mathbf{C}$). For every well-typed decide or install query $(V, q) \in \mathcal{Q}_{\text{sec}}$, compare private names up to consistent renaming and define an equivariant information map f to be exact when

$$\begin{aligned} \text{Exact}(f) &\iff \exists D_f. \forall (V, q) \in \mathcal{Q}_{\text{sec}}, \\ &\quad D_f(\langle f(V), q \rangle_\simeq) = [\text{Ans}(V, q)]_\simeq. \end{aligned}$$

Projection π_J keeps $J \subseteq \{\mathbf{P}, \mathbf{I}, \mathbf{E}, \mathbf{C}\}$ and deletes dependent fields; TargetOnly keeps only caller-proposed target data.

Section A-F gives the exact dependencies and shows that the lower bound concerns information, not storage layout.

b) *Required-outcome preservation:*

$$\text{Preserve}(H, u, \widehat{B}) \iff \forall a \in \text{Req}_u. \exists t \in O_{H,u}, b, \\ \text{Lift}_u(t, a) \wedge (t, b) \in \widehat{B}.$$

Thus every required source outcome appears in an accepted target outcome.

c) *Visible runtime behavior:* $\mathcal{O}_{\text{api}}$ exposes snapshot-bound edit answers, installed edits, protected-call answers, dispatch, settlement, and authenticated returns, while hiding only internal work, Save, stale attempts, and recovered crashes. $\mathcal{B}$ matches ideal and compiled states on AgentSec and these visible fields; section A-L gives the full relation.

Theorem 1 (Exact Edit-Safety). *Clause 1 holds for finite-state $\mathcal{W}$ and finite R . For Clauses 2–5, assume accepted registration, complete control and mediation of every protected call, and crash-stable linearizable transactions within each policy domain. Then they hold for every $u \in \text{Edit}_6 \uplus \{\text{Extend}(\xi)\}$ at every finite reachable AgentSec state.*

- 1) **Faithful registration.** *CheckReg($\mathcal{W}, R$) accepts iff stored λ witnesses $\mathcal{W} \sqsubseteq_\lambda R$. On acceptance, the source workflow’s protected-call traces and $\text{BRuns}(R)$ are order-isomorphic and preserve outcomes, effect identities, new/reuse, the safety answer, and the most permissive monitor; malformed registrations fail before startup.*
- 2) **Exact checking and construction.** *The ideal and compiled machines return the same mutually exclusive and exhaustive answer for $\Theta = \Theta_S(u)$: undefined derivation gives state-preserving Invalid, and $\widehat{B}_\delta^\dagger = \emptyset$ gives state-preserving Reject with no installed realization. Otherwise, for $Y = \text{Compile}(\Theta)$, nonempty $\widehat{B}_\delta^\dagger$, existence of an installed realization, $\text{Ready}(S, u, Y)$, and an AgentSec-preserving installed successor are equivalent. The largest execution family satisfies Preserve, retains recorded evidence for every already-satisfied source outcome, and records every authorized removal.*
- 3) **Each of four state relations is necessary for exact answers.** *Within this field decomposition, the full information view is sufficient, and every projection that discards a relation is insufficient:*

$$\text{Exact}(\pi_J) \iff J = \{\mathbf{P}, \mathbf{I}, \mathbf{E}, \mathbf{C}\}.$$

Every exact decision-and-installation information map distinguishes every separating state–query pair in table II. Maps that erase the call-to-effect link or the link between authorization records and effect identities, as well as TargetOnly, are inexact. Generalized nonblocking can serve as an exact solver only after the trusted controller derives the workflow and outcome-specific completion conditions.

- 4) **Repeated use.** *Every enabled event preserves AgentSec, so the theorem applies after every finite reachable sequence of edits, registered extensions, protected calls, authorization-log updates, installations, and crashes, without a uniform execution-length bound.*

- 5) **Runtime correctness.** *$\mathcal{B}$ is a divergence-insensitive weak bisimulation under obs_{api} , so every good pair satisfies $S_0 \approx_{\mathcal{O}_{\text{api}}}^{\text{di}} I_0$. The observation records snapshot-bound answers, replacement authorization tokens, and returns; time, fairness, and post-Dispatch networking are environmental behavior, and disjoint domains compose asynchronously.*

d) *Proof roadmap:* Registration reflection follows by enumerating the finite protected-call traces and checking their order- and identity-preserving correspondence. Exact checking follows from schema fidelity and the greatest-fixed-point argument: every safe implementation is contained in the computed family, and a nonempty computed family compiles to an implementation. Four explicit pairs of checked states prove the information result; induction over the runtime rules proves repeated use; and a case analysis over edits, protected calls, installation races, returns, and recovery establishes the weak bisimulation. The appendix gives the complete witnesses and case proofs.

VII. MECHANIZATION AND EXECUTABLE VALIDATION

Six Lean modules pass `--trust=0` and form a paper-facing theorem chain. `exact_checker_iff_realization` proves executable admission equivalent to the independently stated declarative realization; `registered_workflow_exact_admission` composes typed registration, exhaustive mediation, and exact admission; and `six_edit_derivation_exact` proves the executable derivation sound and complete for all six edit rules. `exact_edit_installs_trace_safe_monitor` then constructs an atomic installation whose every finite continuation preserves AgentSec, while `fresh_alias_installation_serialized` proves both possible orders between a protected call and installation. `compilation_preload_preserves_agentSec` makes compilation an explicit kernel transition and proves that inactive candidate preloading preserves the same invariant.

All 19 paper-specific tests pass, together with 62 authority and compiler regression tests for supporting components. Across accepted instances with 2–128 outcomes, checker time is 0.11–53.47 ms; rejected instances take 0.11–5.51 ms, and enumerating 6–720 unordered completions takes 0.33–50.35 ms. Section A-M reports the complete measurement protocol, executable coverage, and state-preserving resource policy. The resource-managed wrapper classifies preflight-cap results separately and leaves the authenticated state unchanged.

VIII. RELATED WORK

a) *Recovery, workflow update, and synthesis:* Output-commit recovery reconstructs a fixed computation; generalized nonblocking and Live Synthesis solve supplied transition systems and completion conditions; workflow inheritance and dynamic controller update connect supplied old and new specifications [20], [21], [15], [16], [17], [22], [23], [24]. Our controller derives the edited workflow, protected-effect aliases, still-required outcomes, completion conditions, and replacement

set directly from the authenticated Agent record, then installs the exact most-permissive monitor under fresh and alias races.

b) Agent recovery mechanisms: ACRFence blocks replay after Restore, DART selects checkpoints, Atomix delays settlement, and Rebound makes rollback atomic and auditable [25], [14], [12], [26]. Our theorem subsumes these concerns in one exact criterion for all six derived Agent edits, adds an impossibility certificate and irredundant-state lower bound, and carries the result through atomic installation and arbitrary finite execution.

IX. CONCLUSION

Safe execution editing has an exact criterion: past protected effects remain immutable, and every still-required outcome retains a policy-safe completion. Our checker turns every registered Fork, Restore, Merge, or extension into the unique most-permissive monitor or a finite impossibility certificate. The Exact Edit-Safety theorem unifies exact synthesis, irredundant authenticated state, atomic race serialization, repeated finite execution, and ideal-runtime equivalence, providing a formal foundation for branchable Agent runtimes.

APPENDIX A
SUPPLEMENTARY PROOF APPENDIX

This appendix instantiates the proof objects used by theorem 1 over finite registered call models and atomic policy-domain monitor versions. Independent domains retain the product interpretation stated in the paper.

A. The Declarative Specification and Its Computed Witness

Fix a well-formed safety-check instance $\Theta = (\kappa, H, \Delta, \text{out}, \mathcal{A}, u)$ for which the edit derivation yields $\delta = (\kappa_u, H_u, \mathcal{A}_u, C_u, \eta^-, \eta^+)$. All sets in this subsection are indexed sets, so two equal resolved words belonging to different outcomes remain different elements.

Definition 4 (Required source outcomes). Write $\text{kind}(u)$ for the constructor of u . Call either Fork constructor a Fork and either Restore constructor a Restore. For the matching source row in table I, define

u	$\text{SrcReg}(H, u)$
Fork	$\{(L, C), (R, C)\}$
Restore	$\{\text{cur}, C\}, (k, C_k)\}$
MergeSelect(g, w, σ)	$\{(w, \llbracket T_w \rrbracket), (\bar{w}, \llbracket T_{\bar{w}} \rrbracket)\}$
MergeJoin(g, σ)	$\{(L, \llbracket T_L \rrbracket), (R, \llbracket T_R \rrbracket)\}$

The full set of source outcomes is

$$\text{SourceReq}(H, u) = (\{\text{ctx}\} \times O_{\text{byp}}) \uplus \biguplus_{\substack{(r, D) \\ \in \text{SrcReg}(H, u)}} (\{r\} \times O_D).$$

Write $\text{Done}_H(r, o)$ when the authenticated cursor and progress records residualize outcome o 's pomset to empty. Only RestoreReplace treats completed current outcomes as finished:

$$\text{DoneCur}_H = \{(\text{cur}, o) \mid o \in O_C \wedge \text{Done}_H(\text{cur}, o)\},$$

u	$\text{Completed}_{H, u}$
RestoreReplace(b, k, σ)	DoneCur_H
otherwise	$\emptyset$

Completion observability makes this current-role set either all of $\{\text{cur}\} \times O_C$ or empty. The only sets eligible for signed removal are

$$\text{CurOut} = \{\text{cur}\} \times O_C, \quad \text{Other}_w = \{\bar{w}\} \times O_{\llbracket T_{\bar{w}} \rrbracket},$$

u	$\text{CanRemove}(H, u)$
RestoreReplace(b, k, σ)	$\text{CurOut} \setminus \text{Completed}_{H, u}$
MergeSelect(g, w, σ)	Other_w
otherwise	$\emptyset$

$\text{reqauth}_H(a)$ follows a 's authenticated outcome lineage to the unique signed authority row in G , or to the value-snapshotted copy in K for a checkpoint role. Context identity, carry, and clone preserve that row, while every schema-introduced lineage receives the authority signed with its schema and policy version. Thus reqauth_H is total and immutable on $\text{SourceReq}(H, u)$, including Fork-created L/R copies. For

every $a \in \text{CanRemove}(H, u)$, the removal bundle must contain a signature by $\text{reqauth}_H(a)$ over

$$\begin{aligned} \text{rmmsg}(H, u) = & (\omega, \kappa, \zeta, \sigma, \text{fp}(H), \\ & \text{CanRemove}(H, u), \\ & \text{lineage}(\text{CanRemove}(H, u))). \end{aligned}$$

Without this exact bundle, the edit has no derivation; with it, $\text{Removed}_{H, u} = \text{CanRemove}(H, u)$. The checker requires

$$\begin{aligned} \text{Req}_u = & \text{SourceReq}(H, u) \\ & \setminus (\text{Completed}_{H, u} \uplus \text{Removed}_{H, u}), \\ \text{SourceReq}(H, u) = & \text{Req}_u \uplus \text{Completed}_{H, u} \\ & \uplus \text{Removed}_{H, u}. \end{aligned}$$

Definition 5 (Required-outcome preservation).

$$\begin{aligned} \text{Preserve}(H, u, \hat{B}) \iff & \forall a \in \text{Req}_u. \exists t \in O_{H, u}. b. \\ & \text{Lift}_u(t, a) \wedge (t, b) \in \hat{B}. \end{aligned}$$

Here $\text{Removed}_{H, u}$ contains the typed unselected-arm outcomes for MergeSelect or the current-role outcomes of a pending certified RestoreReplace. $\text{Completed}_{H, u}$ contains the completed current role of RestoreReplace. Both sets are empty for every other row and for a registered extension. Thus $\text{Req}_u = \text{SourceReq}(H, u) \setminus (\text{Completed}_{H, u} \uplus \text{Removed}_{H, u})$, so Preserve covers exactly the source outcomes that are neither completed nor authorized for removal.

Let $\mathfrak{R}_\Theta$ be the finite set of indexed realizations from definition 2, ordered componentwise by

$$(L_1, \hat{B}_1) \sqsubseteq (L_2, \hat{B}_2) \iff L_1 \subseteq L_2 \wedge \hat{B}_1 \subseteq \hat{B}_2.$$

Lemma 8 (Declarative–algorithmic bridge). The declarative predicate $\text{GReal}(\Theta; L, \hat{B})$ has a witness if and only if $\hat{B}_{H, u}^\dagger \neq \emptyset$. When a witness exists, it is unique and equals $(W_{H, u}^\dagger, \hat{B}_{H, u}^\dagger)$. This equality is a theorem relating two independently stated objects, rather than the definition of the ideal machine.

Proof. Let $(L, \hat{B}) \in \mathfrak{R}_\Theta$. Source and target fidelity and policy safety give $\hat{B} \subseteq \hat{B}_0$. For every $(o, b) \in \hat{B}$, every prefix $z \preceq_{\text{pref}} b$ belongs to $L = \text{Pref}(\pi_2 \hat{B})$. Persistent outcome coverage therefore supplies, for every $v \in \text{Compat}(z)$, a pair $(v, b') \in \hat{B}$ extending z . Thus $\hat{B} \subseteq \hat{\Phi}(\hat{B})$. By monotonicity on the finite lattice $2^{\hat{B}_0}$, every post-fixed set is contained in the greatest fixed point, so $\hat{B} \subseteq \hat{B}^\dagger$. Taking projected prefixes yields $L \subseteq W^\dagger$.

Conversely, suppose $\hat{B}^\dagger \neq \emptyset$. The fixed-point equality supplies persistent coverage and, at ϵ , composes with checked lifts to establish Preserve; inclusion in $\hat{B}_0$ supplies target fidelity and safe completed words, prefix closure of $\mathcal{A}$ supplies safety for every generated prefix, and $W^\dagger = \text{Pref}(\pi_2 \hat{B}^\dagger)$ supplies completion nonblocking. Hence $(W^\dagger, \hat{B}^\dagger) \in \mathfrak{R}_\Theta$, and the preceding containment makes it greatest. If two greatest elements existed, each would contain the other componentwise, so they would be equal. If $\hat{B}^\dagger = \emptyset$, the first half places every hypothetical realization inside the empty set, contradicting its required nonemptiness. $\square$

Corollary 1 (Exact compilation boundary). *The ideal edit transition is enabled by the declarative $\text{Spec}(\Theta)$ exactly when the checked compiler returns a nonempty fixed point. On that edge, the ideal language L^* equals the generated language of $\text{Can}(W^\dagger, B^\dagger)$.*

Proof. The first statement is lemma 8. The second combines that lemma with lemma 6. $\square$

Proposition 2 (Output-sensitive compilation). *Fix the unit-cost symbolic-checker model with constant-time finite-map lookup and policy-automaton transitions. Let*

$$\begin{aligned} D &= |\Theta|_{\text{chk}}, & n &= \max_o |X_{p_o}|, \\ \Lambda &= \sum_o \sum_{w \in \text{Lin}(p_o)} (|w| + 1), & N_B &= |\widehat{B}_0|, \\ V_B &= \sum_{(o,b) \in \widehat{B}_0} (|b| + 1). \end{aligned}$$

$$\mathcal{G}_{\text{cov}} = \{((o, b), z, v) \mid (o, b) \in \widehat{B}_0, z \preceq b, v \in \text{Compat}(z)\}.$$

Here D is the authenticated input size, Λ is indexed linearization volume, and $\mathcal{G}_{\text{cov}}$ is the materialized prefix-required-outcome relation. The semantic compiler runs in $O(D + n\Lambda + |\mathcal{G}_{\text{cov}}| + V_B \log(2 + V_B))$ time and $O(D + \Lambda + |\mathcal{G}_{\text{cov}}|)$ space, including its emitted witnesses. Moreover, $N_B \leq \sum_o |\text{Lin}(p_o)|$, $|\mathcal{G}_{\text{cov}}| \leq N_B(n + 1)|O_{C_u}|$, and the number of nonempty strict-removal ranks is at most N_B .

Proof. A backtracking topological-order enumerator emits all indexed linearizations in $O(n\Lambda)$ time, while resolution, policy simulation, and construction of raw and resolved prefix tries are linear in emitted volume. For every support key (z, v) , maintain the number of remaining pairs (v, b') extending z , initially enqueue candidates having a required outcome with zero support, and remove candidates in synchronous batches. When removal makes a support count zero, enqueue exactly the candidates incident to that requirement for the next batch. Each candidate and each incidence in $\mathcal{G}_{\text{cov}}$ is processed at most once, and the batches are exactly $\widehat{B}_i \setminus \widehat{B}_{i+1}$. Store one rank and cause per removed candidate rather than copied sets, and construct $\text{Can}(W, B)$ by bottom-up sorting of finite-trie derivative signatures in $O(V_B \log(2 + V_B))$ time. $\square$

B. Ranked Rejection Certificates

For an empty fixed point, the certificate is

$$C_{\text{fp}} = (\widehat{B}_0, E_0, \dots, E_k), \quad E_i \subseteq \widehat{B}_i \times \text{Pref}(\pi_2 \widehat{B}_i) \times O_{H,u}.$$

The verifier first recomputes the canonical $\widehat{B}_0(\Theta)$ and requires equality with the serialized first component. An entry $((o, b), z, v) \in E_i$ verifies exactly when $z \preceq b$, $v \in \text{Compat}(z)$, and no $(v, b') \in \widehat{B}_i$ extends z . The verifier requires the first projection of E_i to equal $\widehat{B}_i \setminus \widehat{\Phi}(\widehat{B}_i)$, sets $\widehat{B}_{i+1} = \widehat{B}_i \setminus \pi_1(E_i)$, and accepts only when $\widehat{B}_{k+1} = \emptyset$. All sets, prefixes, and compatibility checks are finite and use the canonical order fixed by the checker.

Lemma 9 (Rejection-certificate soundness and completeness). *The verifier accepts C_{fp} iff the descending fixed-point computation reaches $\widehat{B}^\dagger = \emptyset$; in that case no realization exists.*

Proof. The initial equality check binds the certificate to Θ , and every verified round is exactly $\widehat{B}_{i+1} = \widehat{\Phi}(\widehat{B}_i)$, so canonical iteration supplies completeness. For soundness, induction places every post-fixed family inside every $\widehat{B}_i$; an empty final family therefore excludes every nonempty realization. $\square$

C. Exact Certificate Payloads

Let enc be canonical and length-delimited, with κ authenticating $\mathcal{A}$. The symbolic LTS uses domain-separated injective authenticated constructors $\text{Seal}_{\text{prog}}$ and Seal_{cut} , so equal seals have equal payloads. For $\Theta = (\kappa, H, \Delta, \text{out}, \mathcal{A}, u)$, define

$$\begin{aligned} \gamma(\Theta) &= \text{Seal}_{\text{cut}}(\text{enc}(\text{request}, \omega, \zeta, \kappa, H, \Delta, \text{out}, \mathcal{A}, u)), \\ H_{\text{prog}} &= \text{Seal}_{\text{prog}}(\text{enc}(\text{prog}_Z, \widehat{B}_{H,u}^\dagger)). \end{aligned}$$

The source-snapshot seal is

$$\text{cut}_Z = \text{Seal}_{\text{cut}}(\text{enc}(\omega, \zeta, \kappa, H, \Delta, \text{out}, u, \mathcal{C}_{H,u}, H_{\text{prog}}, \eta^-, \eta^+)). \quad (11)$$

Because H contains χ, ν, ρ, β , this seal binds workflow progress, record version, registry state, monitor-entry/version bindings, the authorization log, and the dispatch queue. The verifier re-derives the schema judgment and complete indexed fixed point, reconstructs $(W^\dagger, B^\dagger)$, verifies both seals, and checks the full automaton isomorphism

$$\text{prog}_Z \cong \text{Can}(W^\dagger, B^\dagger),$$

including its initial state, marking, and every labeled transition. Canonical encoding order makes Invalid, Reject, and Installable unique and equivariant. A byte implementation may authenticate enc with signatures or collision-resistant digests. The byte-level guarantee is parameterized by signature unforgeability and digest collision resistance, represented as authenticated identities in the formal model.

D. Supporting Edit and Residual Proofs

1) Schema fidelity:

Full proof of lemma 2. Unique object and schema identifiers select one row and record, while context identity preserves everything outside the hole. The Fork rows certify both clones; RestoreReplace certifies the checkpoint clone and the exact disposition of the source continuation, while RestoreLive certifies the carried source and checkpoint clone. MergeSelect certifies the selected and unselected regions, while MergeJoin certifies both carried regions before its barrier. Clone and carry supply the global bijections and preserve occurrence-outcome membership and lineage-authority rows; constructor induction defines pr_r^u , proves the cover, and supplies every Lift_u witness. The common poststate fixes all remaining fields, preserves or initializes branch cursors, and assigns each schema-introduced lineage its signed authority row, while new identifiers, partial-constructor checks, and lemma 1 preserve uniqueness, acyclicity, and contract well-formedness. Finally, β' and $\mathcal{H}_u^+$ bind every current target occurrence to η^+ . $\square$

2) *Registered extension*: The Agent may name only an immutable record ξ binding the exact predecessor schema, registry, execution-record, and policy roots to finite disjoint additions $D_\Sigma, D_\rho, D_\alpha$ and a successor policy $\mathcal{A}^+$, where D_α contains the signed authority row for every newly introduced outcome lineage. The policy authority and call authorization service authenticate ξ ; existing meanings remain unchanged, new identifiers are fresh and domain-scoped, and $\mathcal{A} \subseteq \mathcal{A}^+$. For fresh η^+ , define

$$\begin{aligned} G^+ &= G \uplus D_\alpha, \\ \Sigma_{\zeta^+} &= \Sigma_\zeta \uplus D_\Sigma, \quad \rho^+ = \rho \uplus D_\rho, \\ H_\xi^+ &= (\omega, G^+, T, K, \Sigma_{\zeta^+}, \rho^+, \beta^+, \\ &\quad \chi, \nu + 1, \eta^+), \\ \mathcal{C}_\xi &= \llbracket T \rrbracket, \\ \text{SourceReq}(H, \xi) &= \text{Req}_\xi = O\llbracket T \rrbracket, \\ \text{Completed}_{H, \xi} &= \emptyset, \quad \text{Removed}_{H, \xi} = \emptyset, \\ \text{Lift}_\xi(t, o) &\iff t = o. \end{aligned}$$

The checked judgment is

$$(\kappa, H, \mathcal{A}) \vdash \text{Extend}(\xi) \Downarrow (\kappa^+, H_\xi^+, \mathcal{A}^+, \mathcal{C}_\xi, \eta, \eta^+).$$

It preserves T, K, χ, Δ , out and every existing row of G, Σ_ζ, ρ , appends D_α to the authenticated execution record, and authorizes no removal; every current outcome, occurrence, effect identity, lineage, authority row, and causal edge therefore remains unchanged. It uses the same checker and atomic installation as an execution edit.

Lemma 10 (Registered extension preserves safety). *For an authenticated $\text{Extend}(\xi)$, the current residual indexed realization is a post-fixed family for the extension instance and is therefore contained in $\widehat{B}_{H, \text{Extend}(\xi)}^\dagger \neq \emptyset$. Consequently, a valid registered extension cannot remove a currently required outcome or leave it without a safe completion.*

3) *Residual coherence*:

Full proof of lemma 5. Because z prefixes a projected member of $\widehat{B}^\dagger$, raw recoverability and streaming in lemma 3 give the recoverability and authorization-log equalities. The indexed residual removes exactly that raw prefix, retains every compatible outcome identity, and gives

$$\text{Compat}_{\mathcal{C}/\text{raw}(z)}(s) = \text{Compat}_{\mathcal{C}}(zs).$$

Together with permission-trace concatenation, streaming resolution yields

$$\widehat{B}_0(\mathcal{C}, \Delta)/z = \widehat{B}_0(\mathcal{C}/\text{raw}(z), \Delta_z),$$

with the same retained outcome identities, policy, and target registry. Thus $\widehat{B}^\dagger/z$ is post-fixed for the residual operator and lies within its greatest fixed point Y . Conversely, streaming and residual policy safety put $zY = \{(o, zy) \mid (o, y) \in Y\}$ inside the original $\widehat{B}_0$. The union $\widehat{B}^\dagger \cup zY$ is post-fixed because prefixes $t \preceq z$ use existing witnesses, while prefixes $t = zs$ use witnesses in Y . Greatestness gives $zY \subseteq \widehat{B}^\dagger$, hence $Y = \widehat{B}^\dagger/z$;

finally, $\text{Pref}(B)/z = \text{Pref}(B/z)$ yields the generated-language equality. $\square$

Proof of lemma 10. An authenticated extension preserves the current workflow exactly, preserves effect-identity resolution and all indexed outcomes, and satisfies $\mathcal{A} \subseteq \mathcal{A}^+$, so every old safe completion and coverage witness remains valid for the successor operator. The current residual family is therefore post-fixed, and greatestness gives its inclusion in a nonempty successor fixed point. $\square$

E. Lean Theorem Map

Six Lean modules machine-check the paper-facing theorem spine. `FiniteCore` proves the executable greatest-fixed-point checker equivalent to the declarative realization predicate, including resolution streaming and greatestness. `RegistrationRefinement` proves typed compilation accepted by registration, complete mediation of registered calls, authenticated identity preservation, and equality between compiled and semantic contracts. `HistoryStructure` proves `deriveEdit_iff`, well-formedness preservation, and executable fixtures for Choice/Parallel Fork, Replace/Live Restore, Select/Join Merge. `OperationalSemantics` proves `kernelTrace_preserves_agentSec`, closure of all six edits, registered-extension safety, and both use-installation race orders. `CompilationBridge` constructs a secure atomic installation from executable admission and reflects every successful installation back to admission. `PaperTheorems` composes these results into exact registered admission, six-edit derivation, compilation-preload preservation, trace-safe installation, and use-installation serialization theorems. The general pomset lift, information lower bounds, and ideal-runtime weak bisimulation complete the theorem in this appendix.

F. Normalized Evidence and Safe Projection

This subsection asks which parts of the authenticated security state are necessary to reproduce every safety-check and installation answer. We first normalize private identifiers so that representation choices cannot distinguish two states, then remove one information factor at a time and construct queries whose correct answers differ.

We make the normalization implicit in `Sec` explicit. Let $\mathcal{Q}_{\text{sec}}$ pair every valid $V = \text{Sec}(S; c, r)$ with a well-sorted $q ::= \text{Decide}(u) \mid \text{TryInstall}(u, Y)$, permitting stale names. `Ans` returns `Compile`($\Theta_V(u)$) for `Decide`, and for `TryInstall` returns the successor bundle exactly when `Ready_V(u, Y)`, otherwise `NoCommit`. An equivariant information map f is exact iff some D_f maps every $\langle f(V), q \rangle_{\cong}$ to $[\text{Ans}(V, q)]_{\cong}$. Let $\mathcal{K}$ be the finite typed key universe at a checked source snapshot, containing tagged outcome, occurrence, effect-identity, authorization record, checkpoint, group, domain, monitor-version, policy-version, and schema-version keys. Public keys comprise policy-authority identifiers and other identifiers declared external by the interface; every other key may be consistently renamed. For a typed query q , $\text{supp}(q)$ is defined recursively over its full syntax, including all identifiers nested inside a supplied

symbolic certificate, monitor, or source-snapshot seal. Let $\mathcal{F}$ be the finite set of typed field paths and assign every semantic base field a unique owner

$$\text{own} : \mathcal{F}_{\text{base}} \longrightarrow \{\mathbf{P}, \mathbf{I}, \mathbf{E}, \mathbf{C}\}.$$

P owns workflow constructor tags, current/checkpoint/target contract nodes, order and outcome membership, outcome-lineage authority rows, schema kind and source fingerprints, lift relations, and authorized-removal declarations. **I** owns occurrence–effect-identity incidence, effect-identity–normalized-request/label/scope rows, occurrence lineage, and logical monitor-entry and authorization-token slots. **E** owns the global resolved trace, current and checkpoint cursors, authorization-record links, retry paths, stored return records, dispatch-queue status, and remaining policy. **C** owns the domain, policy/schema/view versions, private namespace, monitor-version allocation and phase, current monitor-entry ownership, certificate origin and coordinates, and checkpoint snapshot coordinates. These four lists are disjoint and exhaustive for semantic base fields. Certificate bytes, seals, fingerprints, digests, derived permission traces, and foreign-key paths are derived fields rather than additional semantic facts.

Write $f \rightsquigarrow g$ when the value or well-typed interpretation of derived field g depends immediately on field f . The dependency graph is acyclic because canonical encodings are constructed from base relations in a fixed order. Every foreign-key path depends on its target. Authenticated monitor-entry and authorization-token bytes depend on their **I** slot and **C** monitor version and owner. Authorization record reconstruction and the permission trace depend on the **E** authorization record and cursor rows together with their **I** effect identities and labels. Target certificates, source-snapshot seals, and digests depend on every encoded base field in their payload. A normalized full view V is a compatible finite assignment to all base fields together with the unique values of all derivable fields. Compatibility requires equal values on shared typed keys and the well-formedness conditions of definition 1. The natural join $\mathbf{P} \bowtie \mathbf{I} \bowtie \mathbf{E} \bowtie \mathbf{C}$ is the compatible union of these uniquely owned base fields followed by deterministic recomputation of derived fields. It is undefined precisely when shared typed keys disagree or a required foreign-key target is absent.

For $J \subseteq \{\mathbf{P}, \mathbf{I}, \mathbf{E}, \mathbf{C}\}$, start with base fields owned by J and the public sort declarations needed to type the retained records. Repeatedly delete any foreign-key path whose target has been deleted and any derived field having a deleted transitive dependency. The unique fixed point of this deletion is $\text{Ret}(J)$, and $\pi_J(V)$ is V restricted to $\text{Ret}(J)$. This definition prevents an erased fact from surviving inside a seal, digest, archived authorization record-only registry path, or certificate payload.

The incidence-erasure projection $\pi_{\neg oc}$ starts from the full view, deletes the occurrence–effect-identity relation, and applies the same dependency deletion while retaining the occurrence and effect-identity marginals. The projection $\pi_{\neg ar}$ analogously deletes authorization record–effect-identity incidence and every path that reconstructs it, including authorization record–creator,

cursor–authorization record, archived creator–effect-identity paths, and every degree, count, or adjacency summary derived from those paths, while retaining the authorization record, creator, occurrence, and effect-identity marginals. State views are compared modulo a sort-preserving bijection on private keys that fixes all public keys. For possibly different literal queries q, q' , write $(V, q) \cong (V', q')$ when one such bijection maps the full syntax of q to that of q' ; equivalently, $\langle V, q \rangle_{\cong} = \langle V', q' \rangle_{\cong}$.

Lemma 11 (Projection is well defined). *Natural join, dependency deletion, and private renaming commute. Consequently, $[\pi_J(V)]_{\cong}$, $[\pi_{\neg oc}(V)]_{\cong}$, and $[\pi_{\neg ar}(V)]_{\cong}$ do not depend on a representative or on a certificate serialization.*

Proof. A sort-preserving bijection preserves key equality, foreign-key reachability, ownership, and the dependency relation. It therefore maps each deletion round to the corresponding deletion round in the renamed view. Induction on the finite acyclic dependency graph proves that the retained fixed points correspond. Canonical recomposition is equivariant because it is a typed constructor over exactly those retained fields. $\square$

Lemma 12 (Representation-independent observational separation). *Let V_0, V_1 be finite valid full views and let q_0, q_1 be typed queries. If $[\text{Ans}(V_0, q_0)]_{\cong} \neq [\text{Ans}(V_1, q_1)]_{\cong}$, then every exact information map f distinguishes the joint inputs: $(f(V_0), q_0) \not\cong (f(V_1), q_1)$. The claim is representation-independent over arbitrary encodings and characterizes the semantic information factors.*

Proof. Assume toward contradiction that $(f(V_0), q_0) \cong (f(V_1), q_1)$. The joint decoder inputs $\langle f(V_0), q_0 \rangle_{\cong}$ and $\langle f(V_1), q_1 \rangle_{\cong}$ then coincide. Exactness forces equal answer orbits, contradicting the premise. $\square$

G. Bootstrap States and Explicit Separation Pairs

Definition 6 (Checked registration refinement). *A boundary-finite typed Agent workflow $\mathcal{W}$ is a finite-state outcome-labeled LTS whose complete paths end at outcome terminals and whose τ -erased protected-call projections $\partial e = (o, s_1 \cdots s_n)$ form finitely many classes under equality $\equiv_{\partial}$. CheckReg rejects any reachable strongly connected component that can still reach an outcome and contains a protected-call edge, permits erased τ -cycles, and then enumerates the finite classes. R stores a candidate map λ , and $\text{BRuns}(R)$ is the finite outcome-indexed language obtained by linearizing its registered root pomsets and labeling every occurrence by its protected Use. $\mathcal{W} \sqsubseteq_{\lambda} R$ means that λ induces a bijection from boundary classes to $\text{BRuns}(R)$ that preserves outcomes, incidence, order, and each outcome lineage’s signed policy authority. The check also requires effect identities to be exactly the equivalence classes of normalized requests, requires λ to be total, unique, and free of raw instructions, and requires authorization record evolution to commute with Resolve without dropping a required outcome, call site, or authorization record.*

Every separation proof follows the same recipe. Starting from checked initial states, we vary exactly one security-state

Erased	Full-view pair	Exact answers
$u_R(C) := \text{RestoreReplace}(b, k, \sigma)$ with $C_k = C$.		
P	$q_P^i = \text{MergeSelect}(g_i, L, \sigma);$ $T^0 = b_x : X \oplus_{g_0}^{\text{open}} b_y : Y;$ $T^1 = b_x : X \parallel_{g_1} b_y : Y;$ joint inputs are isomorphic after erasure.	V^0 : realize; V^1 : Invalid (wrong mode).
I	$q_I = \text{RestoreLive}(b_L, k, \sigma_I)$ after safe empty-root/ForkChoice/Save; $\Delta = \epsilon, \mathcal{A} = \text{Pref}\{a\};$ $q_{\text{eff}}^0(x) = q_{\text{eff}}^0(y) = d_0;$ $q_{\text{eff}}^1(x) = d_0 \neq d_1 = q_{\text{eff}}^1(y);$ $\ell(d_0) = \ell(d_1) = a.$	$(X \otimes Y_k) \oplus Y$: one new, accept; two new reject <i>aa</i> .
E	$q_E = \text{TryInstall}(u_R(x), Y_0); Y_0$ is compiled at V^0 after safe empty-root/ForkChoice/Save/new; $V^0 \xrightarrow{\text{Dispatch}(d)} V^1$, changing only the release phase.	V^0 : install; V^1 : NoCommit (source snapshot changed).
π_{-ar}	$q_{-ar} = u_R(x)$ after safe empty-root/ForkChoice/Save/new; $q_{\text{eff}}(x) = d, \mathcal{A} = \text{Pref}\{a\}$; same marginals $\{c, x\}, \{p\}, \{d, e\};$ $(q_{\text{eff}}(c), p) = (d, d)$ in $V^0, (e, e)$ in $V^1.$	$\text{reuse}(x_k, d)$: accept; $\text{new}(x_k, d)$: reject <i>aa</i> .
C	$q_C = \text{TryInstall}(u_0, Y_0)$, with full package Y_0 compiled at V^0 , is literal in both; name supplies agree off the monitor-version sort, while current monitor version, ownership, and dependent seals differ.	V^0 : install; V^1 : NoCommit.

TABLE II

FACTOR- AND INCIDENCE-ERASURE PAIRS WITH ISOMORPHIC RETAINED VIEWS AND OPPOSITE EXACT ANSWERS.

factor, keep the projected joint state–query inputs isomorphic, and obtain opposite correct answers; every literal query is fixed except the **P** pair’s private group renaming. The resulting pair proves that an exact abstraction cannot erase that factor.

We construct the finite lower-bound witnesses from one literal empty initial state $S_\emptyset$. A registered call model R compiled from a boundary-finite typed Agent workflow $\mathcal{W}$ contains only finite contracts, schemas, policy, the checked registration map, and authenticated outcome-authority, occurrence, effect-identity, normalized-request, lineage, scope, and label rows. It contains no authorization record, cursor, dispatch-queue entry, monitor version, certificate, return record, monitor state, or installation record. The registration relation has a step $S_\emptyset \xrightarrow{\text{register}(\mathcal{W}, R, n)} \text{Reg}(R, n)$ exactly when $\text{CheckReg}(\mathcal{W}, R)$ accepts and n is a fresh private namespace, followed by checked root installation. Root installation constructs empty authorization record, cursor, and dispatch-queue state, sets $\nu = 0$, activates one monitor version, installs the canonical monitor, and authenticates every current binding and root outcome-authority row.

Lemma 13 (Registration reflection). *For finite-state $\mathcal{W}$ and finite R , $\text{CheckReg}(\mathcal{W}, R)$ accepts iff the registered λ witnesses $\mathcal{W} \sqsubseteq_\lambda R$ in definition 6. On acceptance, the λ -lowerings of source executions modulo $\equiv_\partial$ correspond bijectively to $\text{BRuns}(R)$. The correspondence preserves indexed outcomes, causal order, normalized-request equivalence, authorization record evolution, and the safety-check answer.*

Proof. The SCC check terminates because the state graph is finite. A reachable SCC that can still reach an outcome and con-

tains a protected-call edge can be repeated before that outcome, producing projected words of unbounded length. Conversely, if no such SCC contains a protected-call edge, contracting the SCCs yields a DAG with a finite, enumerable projected language. The checker compares that language with $\text{BRuns}(R)$ and directly checks outcome/call-position incidence, order, effect-identity attestation, total unique registration, exclusion of unregistered calls, and authorization record evolution at every prefix. Therefore it accepts exactly the semantic conjunction defining $\mathcal{W} \sqsubseteq_\lambda R$, rather than an oracle restatement. Boundary projection erases only internal steps, so matching projected paths induce a bijection between the corresponding executions. Normalized-request equality and deterministic Resolve make authorization record evolution commute with this bijection at every prefix. The realization families are therefore order-isomorphic, so the fixed-point operator preserves the safety decision, greatestness, and pullback of the monitor. $\square$

Lemma 14 (Natural bootstrap). *If $\text{CheckReg}(\mathcal{W}, R)$ accepts and R ’s registered root contract and policy have a nonempty greatest realization, registration and checked root installation reach a finite state S_R with $\text{AgentSec}(S_R; c_R, \epsilon)$. The installation record originates at the checked root derivation, while every later replacement records an execution-edit derivation or a registered extension.*

Proof. Registration contributes only the checked R and n , and deterministic allocation constructs every runtime identifier from them. The root judgment and checked lifts establish core/schema coherence, and the recomputed nonempty fixed point establishes required-outcome coherence. Empty runtime progress establishes execution coherence, the canonical program at q_ϵ establishes monitor coherence, and the atomic phase/binding assignments establish monitor-version coherence. Only initial installation creates a root origin, and every successful later installation stores its checked execution-edit or registered-extension derivation. $\square$

The following finite pairs use private names except for each displayed query and permission labels. Each displayed state is reached from $S_\emptyset$ by the stated registration, root-installation, Save, protected-call, and edit steps; omitted fields are identical up to the stated private renaming and dependency deletion.

a) **P**: *workflow structure*: Let X and Y be singleton contracts with distinct occurrences x and y , a shared effect identity, and a universal residual policy. From the same registered singleton root, checked initial installation followed by ForkChoice reaches V_P^0 with private group g_0 , while the corresponding ForkParallel trace reaches V_P^1 with private group g_1 . The two views differ only in

$$V_P^0.T = b_L : X \oplus_{g_0}^{\text{open}} b_R : Y, \quad V_P^1.T = b_L : X \parallel_{g_1} b_R : Y.$$

Use the same registered MergeSelect schema with $q_P^i = \text{MergeSelect}(g_i, L, \sigma)$. The choice view has a unique edit derivation and a nonempty one-step greatest realization, whereas the parallel view has no MergeSelect derivation because its group mode is wrong. Deleting **P** removes the

constructor tag and every authenticator depending on it; the private renaming $g_0 \mapsto g_1$ then makes the retained state–query pairs isomorphic.

b) *Bootstrap convention for the incidence pairs*: Let **1** be the singleton-outcome contract whose unique pomset is empty. Its completion ϵ gives a nonempty greatest realization under $\mathcal{A} = \text{Pref}\{a\}$. Both pairs below install this safe root and then a checked ForkChoice. Each inactive schema row needed to equalize effect-identity marginals is identical in both views and never becomes active.

c) **I**: *occurrence–effect-identity incidence*: Let X, Y be singleton contracts with occurrences x, y , and let d_0, d_1 be effect identities without authorization records, both labeled a . Two checked registered call models have identical topology, schemas, occurrence marginals, and effect-identity marginals but the following incidence:

	$q_{\text{eff}}(x)$	$q_{\text{eff}}(y)$
$V_{\mathbf{I}}^0$	d_0	d_0
$V_{\mathbf{I}}^1$	d_0	d_1

From the installed **1** root, the common Fork schema derives $b_L : X \oplus_g^{\text{open}} b_R : Y$. Each outcome uses one effect identity, so both Forks install under $\text{Pref}\{a\}$. Save b_R before progress and issue the common query $q_{\mathbf{I}} = \text{RestoreLive}(b_L, k, \sigma_{\mathbf{I}})$. Up to clone renaming, the target is $(X \otimes Y_k) \oplus Y$. In $V_{\mathbf{I}}^0$, the left outcome uses one new and one reuse on d_0 , while the bypass outcome uses one new, so every outcome has permission trace a . In $V_{\mathbf{I}}^1$, every left-outcome linearization uses new on both d_0, d_1 , so its word is $aa \notin \mathcal{A}$. The left outcome is compatible with ϵ , so the first $\widehat{\Phi}$ round also removes the otherwise safe bypass completion and leaves the greatest fixed point empty. Erasing occurrence–effect-identity incidence and its dependent rows makes the retained views isomorphic while preserving equal marginals.

d) **E**: *mutable effect progress*: Let C, X be singleton contracts with occurrences c, x , let $q_{\text{eff}}(c) = q_{\text{eff}}(x) = d$, and label d by a . From the safe **1** root, a checked Fork installs $b_L : C \oplus_g^{\text{open}} b_R : X$; save b_R , then use c to complete the left arm, create authorization record p , and prepare its dispatch-queue item. At the resulting $V_{\mathbf{E}}^0$, compile $u_{\mathbf{E}} = \text{RestoreReplace}(b_L, k, \sigma_{\mathbf{E}})$ to the full package Y_0 ; its clone of x aliases d , so Y_0 is installable under $\mathcal{A} = \text{Pref}\{a\}$. Let $V_{\mathbf{E}}^0 \xrightarrow{\text{Dispatch}(d)} V_{\mathbf{E}}^1$, which changes only $\mathbf{E}$'s dispatch-queue release phase and preserves $\mathbf{P}, \mathbf{I}, \mathbf{C}$ literally. For the common literal $q_{\mathbf{E}} = \text{TryInstall}(u_{\mathbf{E}}, Y_0)$, $V_{\mathbf{E}}^0$ installs, whereas $V_{\mathbf{E}}^1$ returns NoCommit because the source-snapshot seal binds the queue phase.

e) $\pi_{\neg ar}$: *authorization record–effect-identity incidence*: For a separate pair, keep the common active row $q_{\text{eff}}(x) = d$, creator c , authorization record p , and effect identities d, e labeled a , but set $q_{\text{eff}}(c) = d$ and $p \mapsto d$ in $V_{\neg ar}^0$, versus $q_{\text{eff}}(c) = e$ and $p \mapsto e$ in $V_{\neg ar}^1$. Both states follow the same safe empty-root/ForkChoice/Save/new shape above, have permission trace a , and use the common query $q_{\neg ar} = \text{RestoreReplace}(b_L, k, \sigma_{\neg ar})$. The cloned x_k is reuse(x_k, d) in $V_{\neg ar}^0$, but new(x_k, d) produces forbidden word aa in $V_{\neg ar}^1$.

The two views have the same creator, authorization record, occurrence, and effect-identity marginals and the same retained incidence $x \mapsto d$. $\pi_{\neg ar}$ deletes the differing archived row for c , the authorization record incidence, and every derived reconstruction path, so the retained views are isomorphic. This proves the necessity of semantic authorization record–effect-identity incidence, not of any particular authorization record–table column.

f) **C**: *installation coordinates*: Let a bootstrap namespace be a family $n = (n_s)_{s \in \text{Sort}}$ of injective, per-sort fresh supplies. From the same literal empty initial state, register the same public call model R with supplies n^0, n^1 that agree on every non-monitor-version sort but mint distinct initial versions $\eta_0 \neq \eta_1$, and perform checked root installation. The resulting $V_{\mathbf{C}}^0, V_{\mathbf{C}}^1$ have literally identical $\mathbf{P}, \mathbf{I}, \mathbf{E}$ base fields, while $\mathbf{C}$'s active monitor version, monitor-entry ownership, certificate coordinates, and dependent authenticators differ. Compile the same registered edit u_0 at $V_{\mathbf{C}}^0$ to obtain $Y_0 = \text{Installable}(\delta_0, W_0, \hat{B}_0, Z_0)$, and submit the same literal query $q_{\mathbf{C}} = \text{TryInstall}(u_0, Y_0)$ to both views. At $V_{\mathbf{C}}^0$, recomputation reproduces cut_{Z_0} , so installation succeeds. At $V_{\mathbf{C}}^1$, recomputation binds η_1 and its owner rather than the η_0 -coordinates sealed in Z_0 , so installation returns NoCommit. Erasing $\mathbf{C}$ recursively erases monitor-version-dependent entry/token rows, certificate coordinates, seals, and authenticators; the identity bijection on every retained sort then makes the retained joint state–query inputs isomorphic.

Proposition 3 (Explicit factor and incidence lower bounds). *Every pair above consists of finite states reached from the common empty initial state and satisfying AgentSec. Each pair has isomorphic projected joint state–query inputs and opposite answer orbits. Therefore every proper factor projection, $\pi_{\neg oc}$, and $\pi_{\neg ar}$ is inexact. Every exact abstraction must distinguish all displayed pairs.*

Proof. Lemma 14 gives the common initial state and valid root installation, while the explicit traces above use only modeled Save, protected-call, and execution-edit transitions. For the incidence pairs, the empty root completes with ϵ , and each installed Fork arm has a one-new completion labeled a . Hence every displayed pre-query state is reachable under $\text{Pref}\{a\}$. Open choice enables the Save steps, and the installed monitors enable each stated left-arm new before that arm becomes a completed, still-addressable continuation. For the **E** pair, Dispatch preserves $\mathbf{P}, \mathbf{I}, \mathbf{C}$ but changes a queue phase bound by Y_0 's source-snapshot seal, giving install versus NoCommit. The displayed schema and resolution calculations give the other opposite answers. Lemma 11 and the specified private-name bijections give joint-input projection equality; for the **P** pair the bijection maps the support of $q_{\mathbf{P}}^0$ to that of $q_{\mathbf{P}}^1$, and for every other pair it fixes the common query support. Any projection omitting at least one factor collapses the corresponding factor pair, while the two incidence erasures collapse their corresponding incidence pairs. Inexactness and the general information lower bound now follow from lemma 12. $\square$

Assigning the same caller-proposed target transition system, completion conditions, and monitor version to the **P** pair makes TargetOnly equal while leaving the authenticated queries and opposite answers unchanged. Thus target-only pre-derivation state is not exact, although generalized nonblocking is exact as a backend after trusted derivation supplies the plant, outcome identities, and conditional markings.

H. Ideal Atomic Execution Machine

An ideal state $I = (\kappa, H, \Delta, \text{out}, \mathcal{A}, c, r)$, with $c = (H_c, u, L^*, \hat{B}^*)$, is exactly the authenticated execution data, current monitor version and entries, installed specification, and resolved prefix. Compilation, monitor representation, transaction phases, authenticators, preloading, and installation microsteps refine this state through the compiled abstraction.

a) *Edit transition*: For current $\Theta_I(u)$, an undefined derivation gives state-preserving Invalid, while a defined derivation with no GReal witness gives state-preserving Reject. Otherwise $\text{Derive}(\Theta_I(u))$ supplies the target and monitor version, and independently defined $\text{Spec}(\Theta_I(u)) = (L^*, \hat{B}^*)$ enables one atomic $\text{Edge}(u, [\mathcal{H}_u^+]_{\cong})$ that installs them and resets r to ϵ . All three answers bind the same current source-snapshot token and are mutually exclusive and exhaustive.

b) *Protected call*: For submitted tool request m , an enabled x passes the ideal guard exactly when its current entry and authorization token bind effect identity d , $\text{canon}_{\kappa}(m) = \text{inv}_{\rho}(d)$, and $ra \in L^*$, where $a = \text{new}(x, d)$ if d has no authorization record and $a = \text{reuse}(x, d)$ otherwise. One atomic transition applies HStep, appends a to χ , advances ν , removes x 's binding, and sets $r := ra$; new also creates the labeled authorization record and queues the normalized request, whereas reuse links to the existing return record. Any failed premise returns $\text{deny}(x)$ without changing state or inventing a return value.

c) *Retry, release, and crash*: Save, authenticated retry, Dispatch, settlement, and recovery after a crash have the same state changes as their compiled rules in table III; compiler computation is absent. This least-generated LTS represents post-Dispatch external-network completion as environmental behavior.

By lemma 4, every reachable partial execution of the ideal machine retains a policy-safe completion until the next edit transition changes the workflow contract. The compiled runtime implements this machine without consulting the full completion set at every call.

I. Compiled State and Atomic Transitions

For $\Theta = (\kappa, H, \Delta, \text{out}, \mathcal{A}, u)$, the checker returns exactly one of

$$\text{CheckOutput}(\Theta) ::= \text{Invalid}([\gamma, C_{\text{str}}]_{\cong}) \mid \text{Reject}([\gamma, C_{\text{fp}}]_{\cong}) \mid \text{Installable}(\delta, W^{\dagger}, \hat{B}^{\dagger}, Z),$$

where

$$Z = (\omega, \zeta, \kappa_Z, \text{cut}_Z, u, H_{\text{prog}}, \eta^-, \eta^+).$$

Invalid carries a verified undefined-derivation transcript, Reject carries the complete ranked elimination certificate, and

Installable carries the unique derivation, coverage witnesses, canonical monitor, and Z . Z authenticates the complete source snapshot, derived target, monitor, and both old and new monitor versions.

The compiled state for one policy domain is

$$S = (\kappa, H, \Delta, \text{out}, \mathcal{A}, \text{phase}, \text{bind}, \text{prog}, q, \text{cert}),$$

where phase records whether each monitor version η is unallocated ($\perp$), inactive, active, or closed. The bootstrap rule runs only after $\text{CheckReg}(\mathcal{W}, R)$ accepts, and its exact construction appears in section A-G. Each installed version stores $c = (\text{src}_c, H_c^+, u, W, \hat{B}, Z)$: the authenticated source snapshot, derived target, edit, safe language, indexed family, and checked certificate.

Definition 7 (Agent security-state invariant). *For $c = (\text{src}_c, H_c^+, u, W, \hat{B}, Z)$ and resolved r , $\text{AgentSec}(S; c, r)$ holds exactly when the following clauses hold.*

- 1) Core and schema coherence. *The state is well formed, $c = \text{cert}(\eta)$, its origin derives the installed target, and every required outcome has one authenticated authority and a checked Lift_u witness.*
- 2) Required-outcome coherence. $\hat{B} = \hat{B}_{H_c, u}^{\dagger} \neq \emptyset$, $W = \text{Pref}(\pi_2 \hat{B})$, and Z verifies both at src_c .
- 3) Execution coherence. $(H.G, H.T) = \text{HRes}((H_c^+.G, H_c^+.T), r)$, $H.\chi = H_c^+.\chi r$, and authorization records, queue entries, return links, and checkpoints match r and advance only monotonically.
- 4) Monitor coherence. $\text{prog}(\eta) \cong \text{Can}(W, \pi_2 \hat{B})$, and $q(\eta)$ is its derivative after r .
- 5) Monitor-version coherence. *Only η is active, $\text{bind}|_{X_{[H], T}} = H.\beta$, and it owns exactly one authenticated entry and token for each current occurrence, while inactive or closed versions expose none.*

a) *Atomic protected call*: For submitted tool request m , set $m^{\circ} = \text{canon}_{\kappa}(m)$. For x bound to monitor entry j in version η , $\text{VerifyToken}(h, \rho(x), j, d, m^{\circ})$ verifies the call authorization service, scope, monitor entry, effect identity, lineage, and signed digest and requires $m^{\circ} = \text{inv}_{\rho}(d)$. Set $a = \text{new}(x, d)$ if $d \notin D_{\Delta}$, otherwise $a = \text{reuse}(x, d)$, and let $\text{Use}_x(m) = \text{Use}(x, j, h, m)$. Its sole state-changing transition is

$$\frac{x \text{ enabled in } T \quad \text{bind}(x) = (j, \eta) \quad \eta = H.\eta \quad \text{phase}(\eta) = \text{active} \quad \text{VerifyToken}(h, \rho(x), j, d, m^{\circ}) \quad q(\eta) \xrightarrow{a}_{\text{prog}(\eta)} q'}{S \xrightarrow{\text{Use}_x(m)/a} S'} \quad (12)$$

The paired update $\text{HStep}((G, T), a)$ removes x , records any choice, preserves the other workflow constructors, and appends a to its branch cursor. $\text{Advance}(S, x, a, q')$ changes exactly

$$\begin{aligned} (H.G, H.T) &:= \text{HStep}((G, T), a), & H.\chi &:= \chi a, \\ H.\nu &:= \nu + 1, & H.\beta &:= \beta \setminus \{x\}, \\ \text{bind} &:= \text{bind} \setminus \{x\}, & q(\eta) &:= q'. \end{aligned}$$

In addition, new appends the immutable labeled authorization record and ordered dispatch-queue item $(d, \text{inv}_\rho(d))$, never an unchecked caller buffer. reuse records the effect identity's stored return link without changing Δ , and all other fields remain unchanged. An authenticated Retry follows the archived link from χ to the authorization record for the same normalized request and returns any stable value without changing state.

b) Answer verification and atomic installation: At its linearization point, the trusted installation service ignores any caller-supplied replacement workflow or monitor version, sets $\Theta_S(u) = (\kappa, H, \Delta, \text{out}, \mathcal{A}, u)$, and re-runs $\text{Derive}(\Theta_S(u))$ from the current state. For a positive derivation, $\text{ReCut}(S, u, Z)$ recomputes the right-hand side of equation (11), while Verify_Z recomputes the fixed point, witnesses, canonical monitor, and both seals. Installation requires

$$\begin{aligned} \kappa &= \kappa_Z, & \text{ReCut}(S, u, Z) &= \text{cut}_Z, \\ \text{Verify}_Z(W_{H,u}^\dagger, B_{H,u}^\dagger, \widehat{B}_{H,u}^\dagger), & & (13) \\ H.\eta &= \eta^-, & \text{phase}(\eta^+) &\in \{\perp, \text{inactive}\}. \end{aligned}$$

$\text{Ready}(S, u, Y)$ holds when $Y = \text{Installable}(\delta, W^\dagger, \widehat{B}^\dagger, Z)$, the current derivation is δ , and all state-dependent premises of equation (13) hold. The committing domain transaction installs $(\kappa_u, H_u, \mathcal{A}_u)$, closes η^- , activates η^+ , binds every target occurrence, installs the canonical program at q^0 , and stores the checked certificate. Only after activation does it expose $\mathcal{H}_u^+$ and emit $\text{Edge}(u, [\mathcal{H}_u^+]_\cong)$; old tokens then fail, while Retry still follows its archived authorization-record link. A stale candidate emits no answer and must be recomputed from the current source snapshot.

J. Event Derivatives and Invariant Preservation

Write $A_1, \dots, A_5$ for the five clauses of AgentSec in their stated order. The following matrix is exhaustive for state-changing transitions inside an installed registered call model. Initial registration occurs before the root safety check; later additive schema and policy growth uses the registered-extension transition listed below.

Lemma 15 (Representative-independent event derivative). *The guarded update induces a partial function ∂ on joint state-event orbits $\langle V, e \rangle_\cong$. If $\text{AgentSec}(S; c, r)$ and $S \xrightarrow{e} S'$, the witness in table III satisfies $\text{AgentSec}(S'; c', r')$ and*

$$[\text{Sec}(S'; c', r')]_\cong = \partial(\langle \text{Sec}(S; c, r), e \rangle_\cong).$$

Proof. All rule guards use typed equality, membership, order, canonical constructors, and authenticated payload equality. These operations are equivariant under a simultaneous private-key renaming of the state and event payload. For an event e , let $G_e(V)$ be its rule guard and $U_e(V)$ its successor update when that guard holds. For every such renaming π , $G_e(V) \iff G_{\pi e}(\pi V)$ and $U_{\pi e}(\pi V) = \pi U_e(V)$. Thus equal pointed orbits have the same definedness and the same successor orbit. This proves representative independence and functionality.

For new and reuse, lemmas 3, 5 and 6 establish the first two rows of the matrix. If x lies at b , HStep appends a to b 's

cursor, while the same atomic protected-call update appends a to the global resolved-occurrence trace χ ; residual associativity gives

$$(\Gamma_b/\text{raw}(\chi_b))/x = \Gamma_b/\text{raw}(\chi_b a).$$

Thus the updated leaf and any later Save remain synchronized with G . For a successful edit transition, lemmas 2, 6 to 8 and 10 establish the third row for all six schema rows and registered extension. Each remaining rule changes only the fields shown in the matrix, and its row checks every invariant clause whose fields change. Thus every successor has the displayed witness and normalized update. $\square$

Corollary 2 (Finite-prefix closure). *Starting from any valid bootstrap, every finite prefix of any sequence of the event classes in table III ends in a state satisfying AgentSec . The sequence length is unbounded, although every individual safety-check instance and event payload is finite. For a finite word $\bar{e}$, the iterated derivative is defined on $\langle V, \bar{e} \rangle_\cong$, where one renaming acts on the initial view and the whole word.*

Proof. The base case is lemma 14, and lemma 15 supplies both the invariant-preserving inductive step and equivariance under one renaming of the remaining event suffix. $\square$

K. Uncontrollable Events and Recovered-State Crashes

The resolved alphabet contains only commits mediated by the protected monitor. Agent request arrival, scheduling, settlement, delivery, and crash are environmental events, but none appends a new or reuse symbol. Request arrival and scheduler choice do not change stored state until a validated transaction commits. Settlement and delivery emit modeled API observations but leave the resolved prefix unchanged. Dispatch is observable only when the oldest queued request is recorded as released, and it also leaves the resolved prefix unchanged. Therefore these uncontrollable events stutter with respect to the conditional-nonblocking plant state. When an environmental action commits a resolved occurrence, the deployment synthesizes the standard supremal controllable sublanguage.

Lemma 16 (Uncontrollable stutter). *If $\text{AgentSec}(S; c, r)$, e is an uncontrollable modeled event, and $S \xrightarrow{e} S'$, then the resolved prefix and canonical monitor state are unchanged and $\text{AgentSec}(S'; c, r)$ holds.*

Proof. Request arrival and scheduling do not change stored state, while denial, retry, retrieval, delivery, and failure recovery preserve the recovered security state. Dispatch and settlement change only the dispatch queue, authorization-record status, or stored return record permitted by execution coherence. None changes $H.T$, $H.\chi$, the permission trace derived from Δ , r , or $q(\eta)$, and the corresponding row of table III preserves every other invariant clause. $\square$

The crash theorem instantiates the recovered-state interface of section III with linearizable, crash-stable domain transactions. Let D be the stored state, v the unfinished in-memory state, and $\bar{t} = t_1 \cdots t_n$ the finite set of domain transactions overlapping a failure, ordered by a linearizability witness. Linearizability

Event class	State update	Witness	Preservation argument
new use	Apply HStep to (G, T) ; append $\text{new}(x, d)$ to global χ ; advance ν, q ; append one authorization record and normalized-request queue item; remove the current binding.	$(c, r\text{new}(x, d))$	A_1 : registry and schema unchanged. A_2 : residual coherence. A_3 : unique authorization record and ordered preparation. A_4 : canonical derivative. A_5 : remove exactly x 's binding.
reuse use	Apply HStep to (G, T) ; append $\text{reuse}(x, d)$ to global χ ; advance ν, q ; add only the return-record link; remove current binding.	$(c, r\text{reuse}(x, d))$	A_1 : authenticated effect identity retained. A_2 : residual coherence. A_3 : link names an earlier authorization record and the permission trace is unchanged. A_4 : canonical derivative. A_5 : remove exactly x 's binding.
Successful edit transition, six edits or registered extension	Install derived $(\kappa^+, H^+, \mathcal{A}^+)$, checked monitor and certificate; close η^- ; activate η^+ ; publish new bindings.	(c_δ^+, ϵ)	A_1 : schema fidelity and lifts. A_2 : declarative-algorithmic bridge and checker. A_3 : source authorization records and dispatch queue retained. A_4 : install at q_ϵ . A_5 : complete monitor replacement.
Verified Invalid or Reject	Emit the certificate for the current checked state; no stored-state change.	(c, r)	All clauses unchanged; γ binds the visible answer to this source snapshot.
Save	Snapshot the current residual/cursor pair in an authenticated checkpoint; advance ν .	(c, r)	A_1 : immutable well-typed extension. A_2, A_4, A_5 : unchanged. A_3 : exactly the permitted K extension.
Preload	Allocate or replace content only in an inactive unbound monitor version.	(c, r)	A_1 – A_4 : unchanged. A_5 : inactive and unbound remain true; no authorization token is exposed.
Retry, retrieval, delivery	Emit $\text{Return}(t, [v]_\cong)$ for an authenticated stored value; no security-state change.	(c, r)	All clauses unchanged; retry additionally checks the same normalized request and authorization record link.
Denial	Return $\text{deny}(x)$; no state mutation.	(c, r)	All clauses unchanged.
Dispatch(d)	Mark exactly the oldest prepared dispatch-queue item released.	(c, r)	A_3 : release order remains a prefix of preparation order. Other clauses unchanged.
Settlement	Emit $\text{Settle}(d, [v]_\cong)$, advance one authorization record phase, and fill its stable return record.	(c, r)	A_3 : effect identity, normalized request, creator, permission label, and order are immutable. Other clauses unchanged.
Stale or malformed candidate	No stored-state change.	(c, r)	All clauses unchanged; the attempt is τ and emits no verdict.
Crash and recovery	Discard unfinished work and expose the same recovered stored state.	(c, r)	All clauses unchanged under the assumed crash-stable transaction boundary.

TABLE III
EXHAUSTIVE EVENT CLASSES AND PRESERVATION OF THE FIVE AgentSec CLAUSES.

and crash stability select a prefix $\bar{t}_{\leq k}$ containing exactly the transactions that linearized before the crash and provide

$$\text{Recover}(D, v, \bar{t}) = \text{commit}_{t_k} \circ \dots \circ \text{commit}_{t_1}(D),$$

with the empty prefix interpreted as D . No recovered image contains a strict subset of one modeled atomic update. In particular, a new transaction cannot expose a authorization record without its global-trace record, branch-local cursor record, and dispatch-queue preparation. Installation cannot activate a new monitor version without closing the old version and installing all target bindings.

Lemma 17 (Recovered-state crash refinement). *Let $S_0 \rightarrow S_1 \rightarrow \dots \rightarrow S_n$ be the modeled path whose guarded rule transitions correspond to $t_1, \dots, t_n$ in the linearizability order above. If stored state D refines S_0 , then recovery after the selected committed prefix $t_1 \dots t_k$ refines S_k . At recovered-state LTS boundaries, crash is consequently a τ self-loop.*

Proof. Induction on k starts from D refining S_0 . At each step, the complete atomic update commit_{t_i} is the guarded rule update in the corresponding row of table III, so it carries a concrete image refining S_{i-1} to one refining S_i . The recovery equation applies exactly the first k such updates and no partial update, yielding the claim. Once the LTS records the recovered image, crashing again changes only volatile state and hence becomes a self-loop. $\square$

L. Agent-API Weak Bisimulation After Recovery

We now prove the runtime-correctness clause of theorem 1 by exhausting the observable and silent cases. For $K = (\kappa, H, \Delta, \text{out}, \mathcal{A})$, $\text{Current}(H) = X_{[H, T]}$, and $I = (K, c_I, r_I)$, let $\alpha_I(I) = (K, c_I, r_I, H.\eta, H.\beta)$ and $\alpha_S(S; c, r) =$

$(K, \text{icut}(c), r, H.\eta, \text{bind}|_{\text{Current}(H)})$. Define IBS iff some c, r satisfy $\text{AgentSec}(S; c, r)$ and $\alpha_I(I) = \alpha_S(S; c, r)$; write $\xRightarrow{\ell} = \tau^* \ell \tau^*$ and $\xRightarrow{\tau} = \tau^*$. Fix IBS with witness (c, r) . Equality of $\alpha_I(I)$ and $\alpha_S(S; c, r)$ gives both machines the same κ , current policy, normalized execution record, authorization log, dispatch queue, stored return records, ideal projection, current partial execution, monitor version, and monitor-entry bindings.

a) *Edit transitions:* The normalization equality makes both machines form the same current $\Theta_I(u)$ and source-snapshot token $\gamma(\Theta_I(u))$. If derivation is undefined, deterministic failed-check reflection makes both sides emit the same canonical Invalid orbit and take a state self-loop. If derivation exists but no declarative realization exists, lemma 8 and deterministic fixed-point checking make both sides emit the same canonical Reject orbit and take a state self-loop. Otherwise the ideal machine enables its edit transition and, by lemma 8, the compiled checker accepts exactly the same edit and installs the same largest language. Compilation takes an explicit Preload transition in the compiled LTS: it records the candidate key in inactive metadata, exposes no authorization token or verdict, and preserves AgentSec. The ideal side takes its single atomic edge, and both successors normalize to the same $(\kappa_u, H_u, \mathcal{A}_u)$, ideal projection, empty prefix, active monitor version, and target bindings. Deterministic allocation gives the same authorization-token bundle up to private renaming, so both expose $\text{Edge}(u, [\mathcal{H}_u^+]_\cong)$. Conversely, every successful compiled installation has passed the schema, fixed-point, certificate, and source-snapshot checks, so lemma 8 enables the matching ideal edge. This argument applies uniformly to all six execution edits and the registered extension.

b) *Protected calls and submitted requests*: For a submitted tool request m , both guards compute the same $m^\circ = \text{canon}_\kappa(m)$. The compiled authorization-token check verifies the call authorization service, scope, occurrence, monitor entry, effect identity, lineage, signed digest, current monitor version, and equality $m^\circ = \text{inv}_\rho(d)$. The ideal guard requires the same normalized binding and request equality. Both sides then inspect the same set D_Δ of effect identities with authorization records, so they choose the same $a \in \{\text{new}(x, d), \text{reuse}(x, d)\}$. By corollary 1 and lemma 6, ra is enabled by the compiled derivative exactly when $ra \in L^*$ in the ideal machine.

In the new case, both successors apply the same paired residual, append the same symbol to the global trace and branch-local cursor, advance ν , create the same effect-identity-bound authorization record, and prepare the same normalized request. The compiled dispatch queue stores $(d, \text{inv}_\rho(d))$, never the caller buffer, while the ideal queue stores (d, m°) ; the verified equality makes these records identical. In the reuse case, both successors apply the same paired update, append the same symbol that creates no authorization record to the global record and branch cursor, and bind the occurrence to the same stored return record without changing Δ or the dispatch queue. Both successors remove x 's current binding; because the remaining workflow no longer contains x , their normalized current-binding components remain equal. In either case, lemma 15 re-establishes $\mathcal{B}$ with witness (c, ra) .

If any occurrence, monitor entry, authorization token, normalized request, monitor version, branch, or monitor transition check fails, both normalized guards fail. Both machines emit $\text{deny}(x)$, preserve state, and remain related.

c) *Use–installation races*: Use and installation serialize on the same execution and monitor-version records. If a new use commits first, it changes $(G, T, \Delta, \chi, \nu, \text{out})$, all of which are bound directly or through H by the source-snapshot seal, so the prepared installation becomes stale. If an reuse use commits first, it changes (G, T, χ, ν) even though Δ and the dispatch queue are unchanged, so the prepared installation again becomes stale. If installation commits first, it atomically closes η^- , activates η^+ , refreshes every current monitor entry and authorization token, and exposes the replacement authorization tokens only after activation. The old new or reuse use then fails the current-monitor-version guard and is denied without progress. The ideal atomic order has exactly the same two cases, so neither side admits a third race outcome.

d) *Replies, releases, and silent steps*: An authenticated retry, authorization-token retrieval, or result delivery leaves the security state unchanged and follows the same authenticated link on both sides to expose $\text{Return}(t, [v]_\cong)$. $\text{Dispatch}(d)$ is enabled at the same dispatch-queue head and makes the same observable release update. $\text{Settle}(d, [v]_\cong)$ is enabled for the same authorization record and produces the same stable result and normalized update. Save takes matching τ -steps. Compiled preload and stale or malformed candidate attempts are implementation-only τ -steps matched by zero ideal steps because α_S omits their inactive contents and they emit no verdict. Crash is matched by a τ self-loop on each recovered-

state machine by lemma 17.

Theorem 2 (Agent-API weak bisimulation after recovery, detailed). $\mathcal{B}$ is a divergence-insensitive weak bisimulation under obs_{api} .

Proof. The least-generated rule sets in sections A-H and V and table III make the preceding cases exhaustive. For every step from either side, the corresponding case constructs a path with observation $\text{obs}_{\text{api}}(\ell)$ and a successor related by the witness in that matrix. The construction is symmetric because successful ideal guards and successful compiled guards were shown equivalent, while implementation-only steps preserve the common abstraction. Arbitrary repetition of silent preload, stale-candidate, Save , or crash steps is ignored, which is precisely divergence-insensitive weak matching. $\square$

M. Executable-Artifact Evaluation

The performance script runs every sample in a fresh CPython 3.12 process and reports the median of five repetitions on a 16-vCPU AMD EPYC 7R12. The 19 paper-specific tests exercise indexed pomset resolution, the greatest fixed point, and finite fixtures for all six constructors. The executable model covers the checker core and all six constructor fixtures, and the theorem development establishes immutable schema refinement, authority-bound removal certificates, and the complete domain-wide monitor-version protocol. Above 128 outcomes, eight occurrences per outcome, 50,000 linearizations per outcome, or 100,000 indexed completions, the executable wrapper returns $\text{ResourceLimit}(b, n)$ before search and performs no installation. The wrapper classifies this state-preserving preflight result separately from the semantic verdicts of the uncapped finite compiler used by the theorem.

AI USE ACKNOWLEDGMENT

OpenAI Codex was used throughout the manuscript for initial prose drafting and revision, and in the executable artifact for code scaffolding and test generation. It also served as an interactive aid for literature discovery, counterexample search, and the discussion and refinement of the formal theory, including definitions, assumptions, theorem statements, and proof structure. The authors chose and directed the research, made all final theoretical and methodological decisions, independently verified the cited sources, theorem statements, proofs, code, and experimental results, and remain responsible for every claim, proof, result, and citation.

REFERENCES

- [1] OpenAI, “Codex App Server,” Official documentation, 2026, accessed 2026-07-31. [Online]. Available: <https://learn.chatgpt.com/docs/app-server.md>
- [2] Anthropic, “Manage sessions,” Official documentation, 2026, accessed 2026-07-31. [Online]. Available: <https://code.claude.com/docs/en/sessions>
- [3] —, “Checkpointing,” Official documentation, 2026, accessed 2026-07-31. [Online]. Available: <https://code.claude.com/docs/en/checkpointing>
- [4] C. Wang and Y. Zheng, “Fork, explore, commit: Os primitives for agentic exploration,” 2026, agentic OS Workshop 2026. [Online]. Available: <https://arxiv.org/abs/2602.08199>
- [5] T. Wu, C. Chang, L. Cao, W. Gao, and W. Wang, “Crab: A semantics-aware checkpoint/restore runtime for agent sandboxes,” 2026. [Online]. Available: <https://arxiv.org/abs/2604.28138>
- [6] Y. Zhang, “Agent libos: A runtime substrate for capability-controlled self-evolving llm agents,” 2026. [Online]. Available: <https://arxiv.org/abs/2606.03895>
- [7] S. Yu, D. Chong, A. Nandi, D. Soylu, J. Sun, C. D. Manning, and W. Shi, “Shepherd: Enabling programmable meta-agents via reversible agentic execution traces,” 2026. [Online]. Available: <https://arxiv.org/abs/2511.10913>
- [8] Y. Li, S. Ping, X. Chen, X. Qi, Z. Wang, Y. Luo, and X. Zhang, “AgentGit: A version control framework for reliable and scalable LLM-powered multi-agent systems,” 2025. [Online]. Available: <https://arxiv.org/abs/2511.00628>
- [9] LangChain AI, “Use time-travel,” 2026, accessed 2026-08-03. [Online]. Available: <https://docs.langchain.com/oss/python/langgraph/use-time-travel>
- [10] S. G. Patil, T. Zhang, V. Fang, N. C., R. Huang, A. Hao, M. Casado, J. E. Gonzalez, R. A. Popa, and I. Stoica, “GoEX: Perspectives and designs towards a runtime for autonomous LLM applications,” 2024. [Online]. Available: <https://arxiv.org/abs/2404.06921>
- [11] E. Y. Chang and L. Geng, “SagaLLM: Context management, validation, and transaction guarantees for multi-agent LLM planning,” *Proceedings of the VLDB Endowment*, vol. 18, no. 12, pp. 4874–4886, 2025. [Online]. Available: <https://www.vldb.org/pvldb/vol18/p4874-chang.pdf>
- [12] B. Mohammadi, N. Potamitis, L. Klein, A. Arora, and L. Bindshaedler, “Atomix: Timely, transactional tool use for reliable agentic workflows,” 2026. [Online]. Available: <https://arxiv.org/abs/2602.14849>
- [13] Z. Chen, H. Liu, D. Xu, D. Dong, J. Li, B. Pu, and J. Zhai, “Cordon: Semantic transactions for tool-using llm agents,” 2026. [Online]. Available: <https://arxiv.org/abs/2606.17573>
- [14] K. Yang, P. Li, Z. Wu, K. Xu, H. Huang, and X. Huang, “DART: Semantic recoverability for structured tool agents,” 2026. [Online]. Available: <https://arxiv.org/abs/2605.23311>
- [15] M. H. de Queiroz, J. E. R. Cury, and W. M. Wonham, “Multitasking supervisory control of discrete-event systems,” *Discrete Event Dynamic Systems*, vol. 15, no. 4, pp. 375–395, 2005.
- [16] R. Malik and R. Leduc, “Generalised nonblocking,” in *Proceedings of the 9th International Workshop on Discrete Event Systems*, 2008, pp. 340–345.
- [17] B. Finkbeiner, F. Klein, and N. Metzger, “Live synthesis,” *Innovations in Systems and Software Engineering*, vol. 18, no. 3, pp. 443–454, 2022. [Online]. Available: <https://doi.org/10.1007/s11334-022-00447-5>
- [18] L. Aceto, I. Cassar, A. Francalanza, and A. Ingólfssdóttir, “On runtime enforcement via suppressions,” in *29th International Conference on Concurrency Theory (CONCUR 2018)*, ser. Leibniz International Proceedings in Informatics, vol. 118. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2018, pp. 34:1–34:17.
- [19] J. A. Brzozowski, “Derivatives of regular expressions,” *Journal of the ACM*, vol. 11, no. 4, pp. 481–494, 1964.
- [20] R. E. Strom and S. Yemini, “Optimistic recovery in distributed systems,” *ACM Transactions on Computer Systems*, vol. 3, no. 3, pp. 204–226, 1985.
- [21] E. Y. Elnozahy, L. Alvisi, Y.-M. Wang, and D. B. Johnson, “A survey of rollback-recovery protocols in message-passing systems,” *ACM Computing Surveys*, vol. 34, no. 3, pp. 375–408, 2002.
- [22] W. M. P. van der Aalst and T. Basten, “Inheritance of workflows: An approach to tackling problems related to change,” *Theoretical Computer Science*, vol. 270, no. 1–2, pp. 125–203, 2002.
- [23] L. Nahabedian, V. A. Braberman, N. D’Ippolito, S. Honiden, J. Kramer, K. Tei, and S. Uchitel, “Dynamic update of discrete event controllers,” *IEEE Transactions on Software Engineering*, vol. 46, no. 11, pp. 1220–1240, 2020.
- [24] G. Amram, S. Maoz, I. Segall, and M. Yossef, “Dynamic update for synthesized GR(1) controllers,” in *Proceedings of the 44th International Conference on Software Engineering (ICSE)*. ACM, 2022, pp. 786–797.
- [25] Y. Zheng, Y. Yang, W. Zhang, and A. Quinn, “ACRFence: Preventing semantic rollback attacks in agent checkpoint-restore,” 2026, coDAIM Workshop 2026. [Online]. Available: <https://arxiv.org/abs/2603.20625>
- [26] Q. Burke, A. Vahldiek-Oberwagner, M. Swift, and P. McDaniel, “It’s a feature, not a bug: Secure and auditable state rollback for confidential cloud applications,” 2025, accepted at the 2026 IEEE Symposium on Security and Privacy. [Online]. Available: <https://arxiv.org/abs/2511.13641>